

OpenMP Tutorial

Performance: Vectorization

Christian Terboven

Michael Klemm

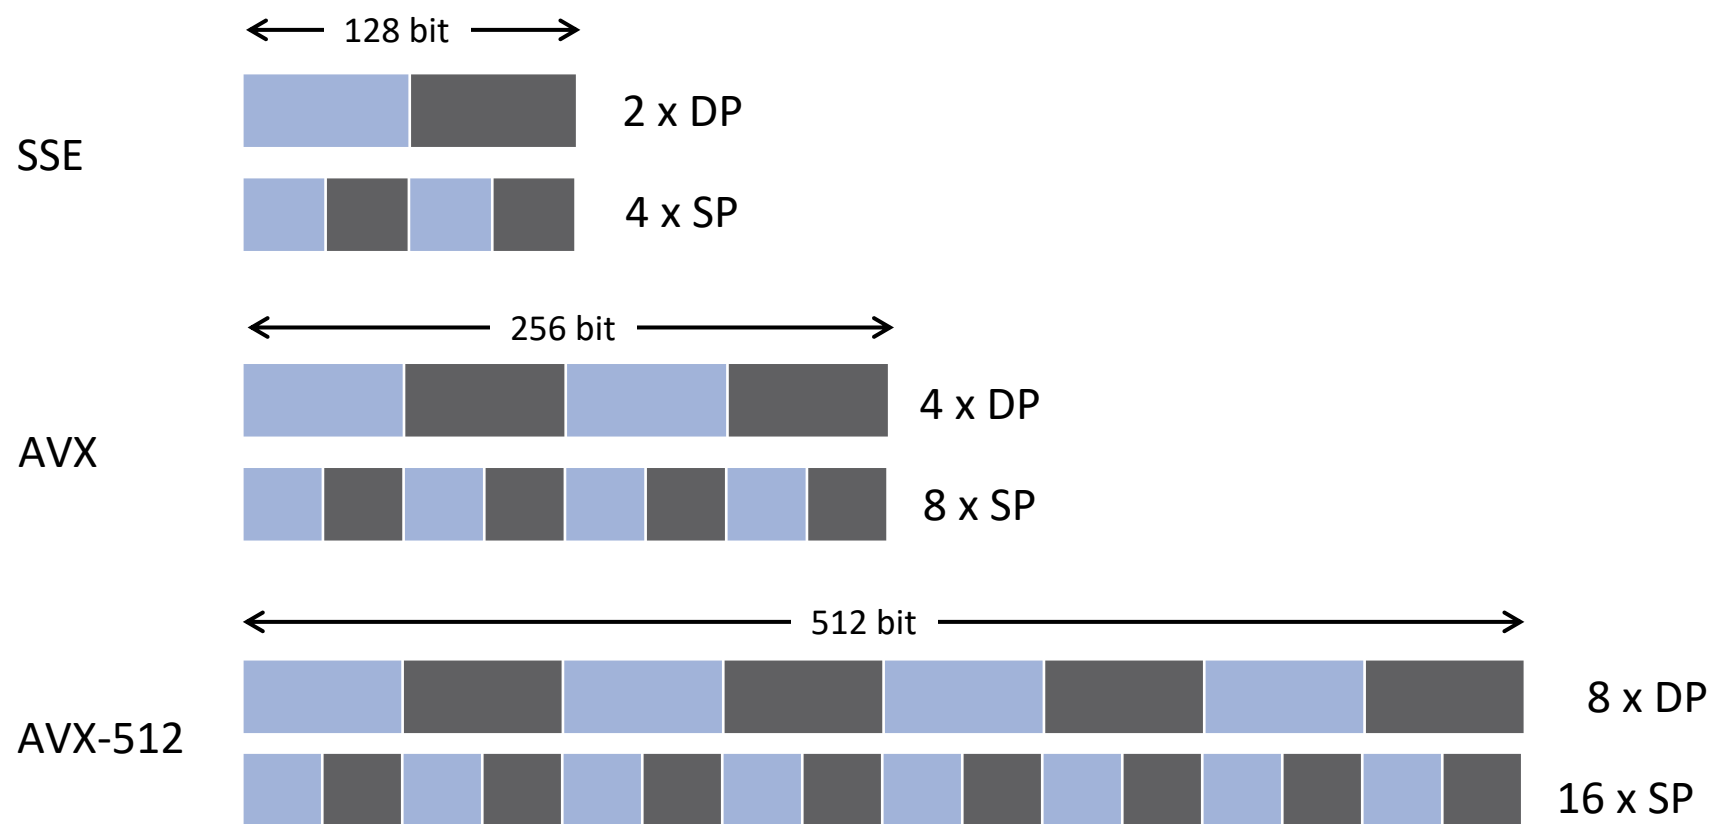
The logo for RWTH Aachen University, with "RWTHAACHEN" in blue and "UNIVERSITY" in a lighter blue below it.The OpenMP logo, featuring the text "OpenMP" in a teal color with a horizontal line underneath the "Open" part.

Topics

- Exploiting SIMD parallelism with OpenMP
- Using SIMD directives with loops
- Creating SIMD functions

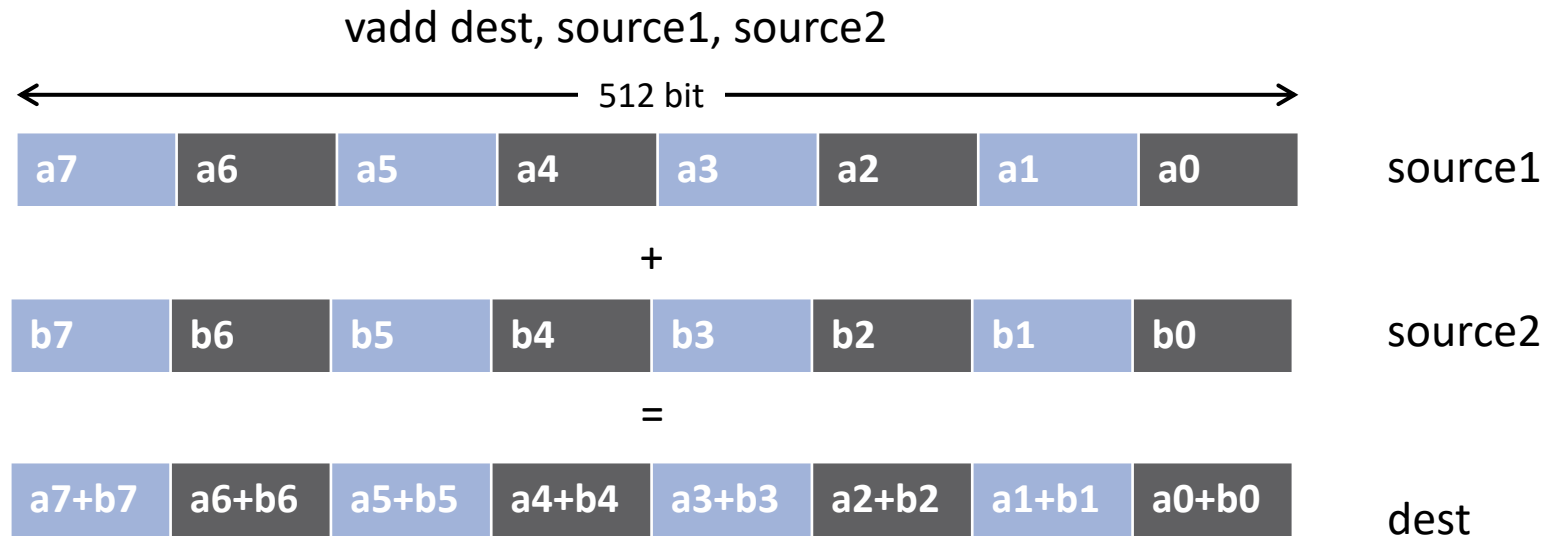
SIMD on x86 Architectures

- Width of SIMD registers has been growing in the past:



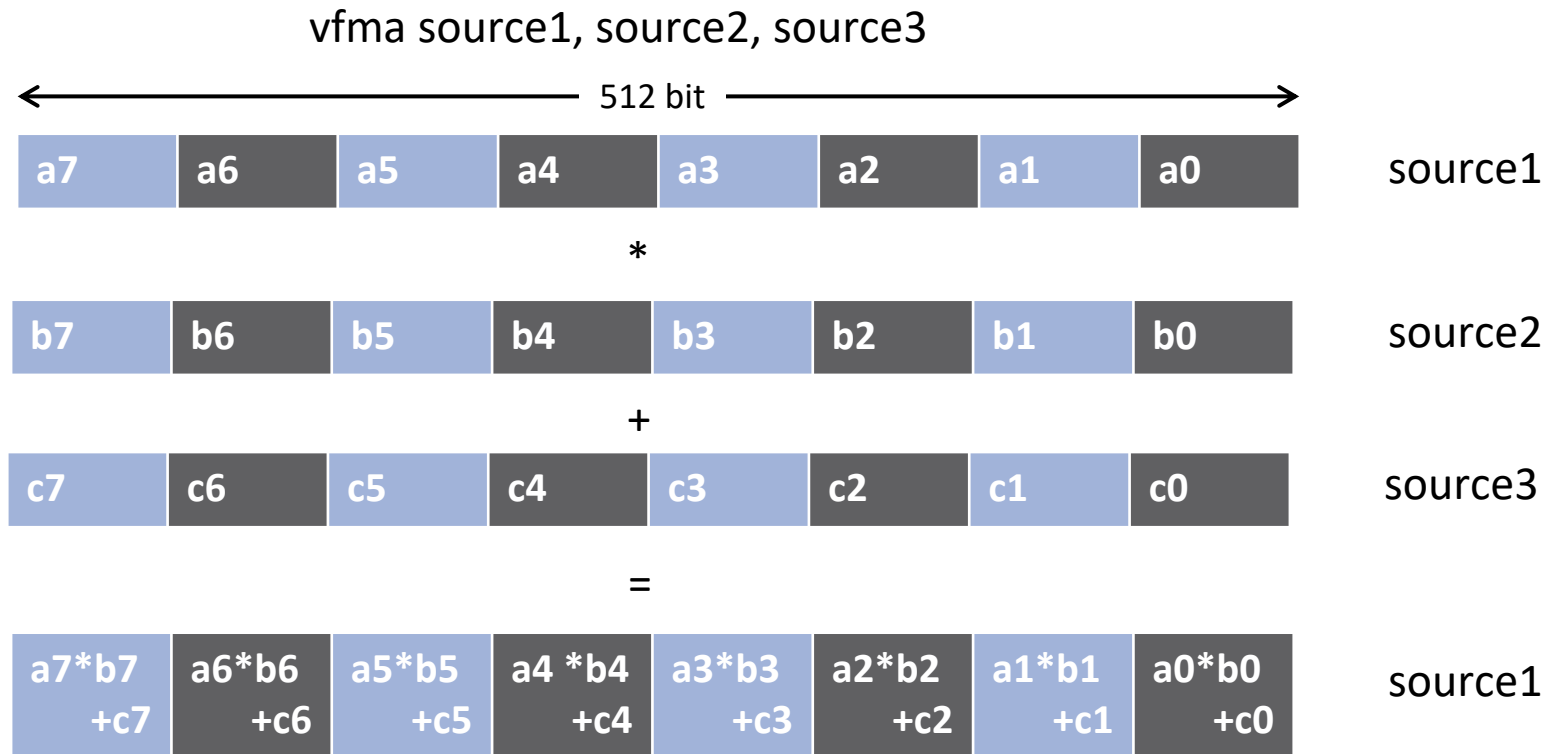
More Powerful SIMD Units

- SIMD instructions become more powerful



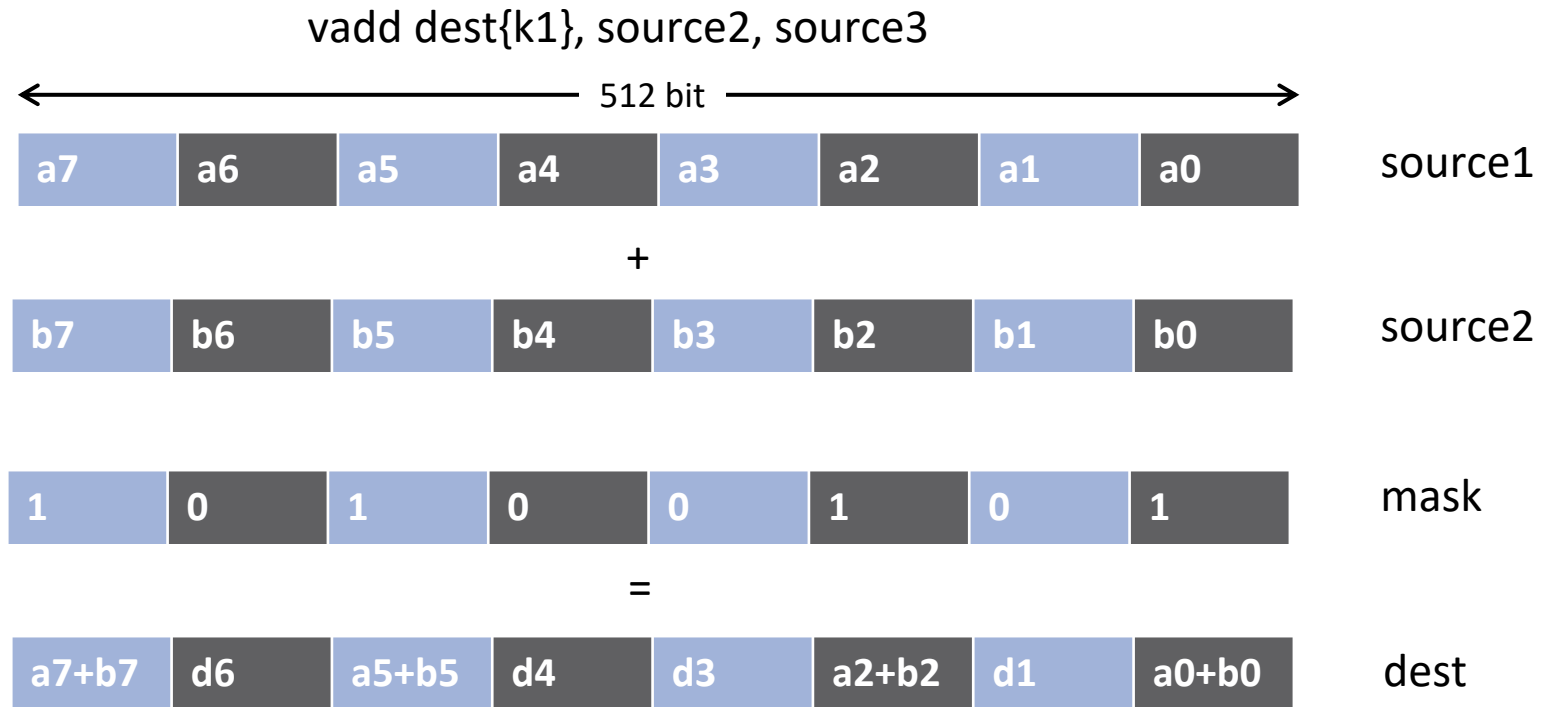
More Powerful SIMD Units

- SIMD instructions become more powerful



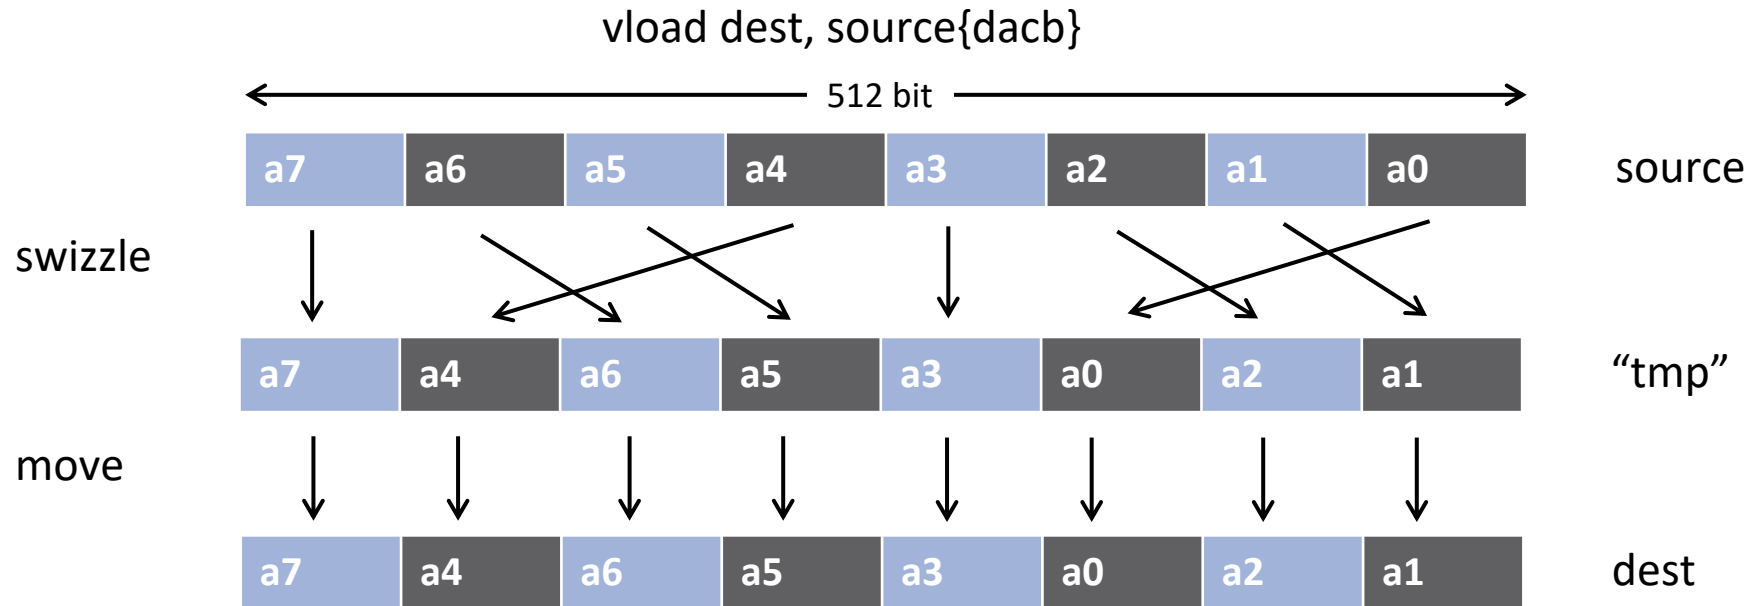
More Powerful SIMD Units

- SIMD instructions become more powerful



More Powerful SIMD Units

- SIMD instructions become more powerful



Auto-vectorization

■ Compilers offer auto-vectorization as an optimization pass

→ Usually, part of the general loop optimization passes

→ Code analysis detects code properties that inhibit SIMD vectorization

?

→ Heuristics determine if SIMD execution might be beneficial

→ If all goes well, the compiler will generate SIMD instructions

■ Example: clang/LLVM

→ -fvectorize

→ -Rpass=loop-.*

→ -mprefer-vector-width=<width>

GCC

-ftree-vectorize

-ftree-loop-vectorize

-fopt-info-vec-all

Intel Compiler

-vec (enabled w/ -O2)

-qopt-report=vec

Why Auto-vectorizers Fail

■ Data dependencies

■ Other potential reasons

- Alignment
- Function calls in loop block
- Complex control flow / conditional branches
- Loop not “countable”
 - e.g., upper bound not a runtime constant
- Mixed data types
- Non-unit stride between elements
- Loop body too complex (register pressure)
- Vectorization seems inefficient

■ Many more ... but less likely to occur

Data Dependencies

- Suppose two statements S1 and S2
- S2 depends on S1, iff S1 must execute before S2
 - Control-flow dependence
 - Data dependence
 - Dependencies can be carried over between loop iterations
- Important flavors of data dependencies

FLOW

s1: a = 40

b = 21

s2: c = a + 2



ANTI

b = 40

s1: a = b + 1

s2: b = 21



Loop-Carried Dependencies

- Dependencies may occur across loop iterations

→ Loop-carried dependency

- The following code contains such a dependency:

```
void lcd_ex(float* a, float* b, size_t n, float c1, float c2)
{
    size_t i;
    for (i = 0; i < n; i++) {
        a[i] = c1 * a[i + 17] + c2 * b[i];
    }
}
```

Loop-carried dependency for $a[i]$ and $a[i+17]$; distance is 17.

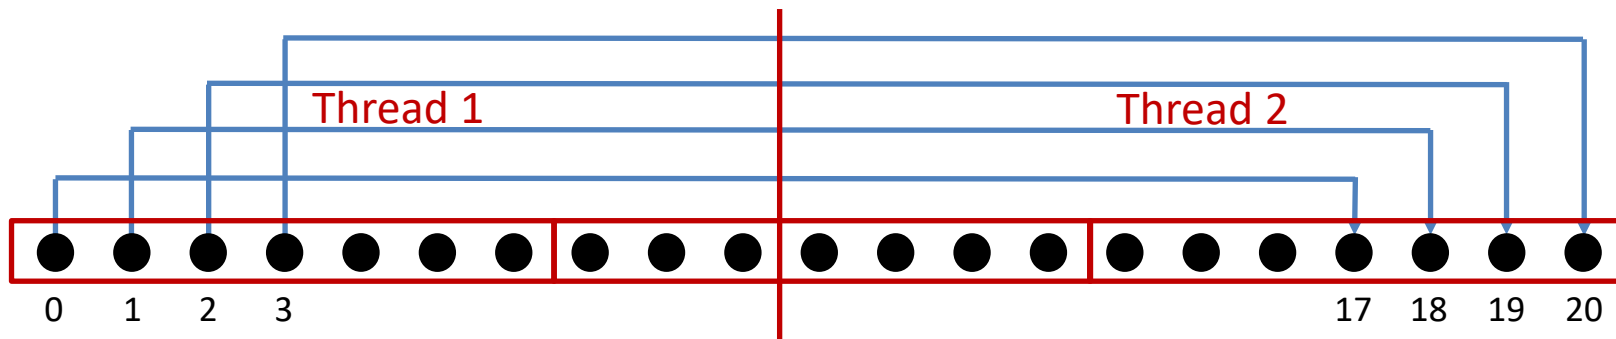
- Some iterations of the loop have to complete before the next iteration can run

→ Simple trick: Can you reverse the loop w/o getting wrong results?

Loop-carried Dependencies

■ Can we parallelize or vectorize the loop?

```
void lcd_ex(float* a, float* b, size_t n, float c1, float c2) {  
    for (int i = 0; i < n; i++) {  
        a[i] = c1 * a[i + 17] + c2 * b[i];  
    }  
}
```



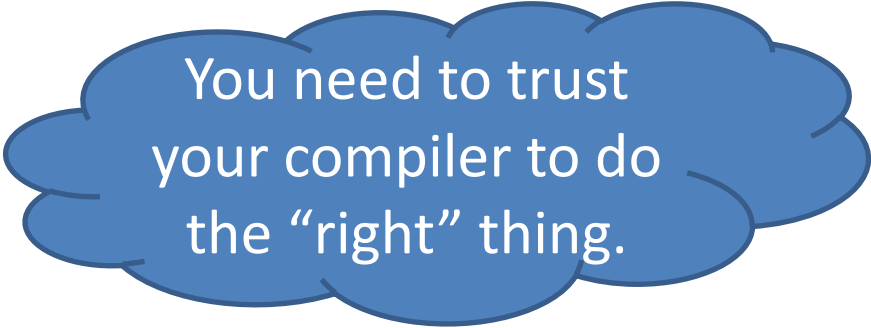
- Parallelization: no
(except for very specific loop schedules)
- Vectorization: yes
(iff vector length is shorter than any distance of any dependency)

In a Time Before OpenMP 4.0

■ Support required vendor-specific extensions

- Programming models (e.g., Intel® Cilk Plus)
- Compiler pragmas (e.g., `#pragma vector`)
- Low-level constructs (e.g., `_mm_add_pd()`)

```
#pragma omp parallel for
#pragma vector always
#pragma ivdep
for (int i = 0; i < N; i++) {
    a[i] = b[i] + ...;
}
```



You need to trust
your compiler to do
the “right” thing.

SIMD Loop Construct

■ Vectorize a loop nest

- Cut loop into chunks that fit a SIMD vector register
- No parallelization of the loop body

■ Syntax (C/C++)

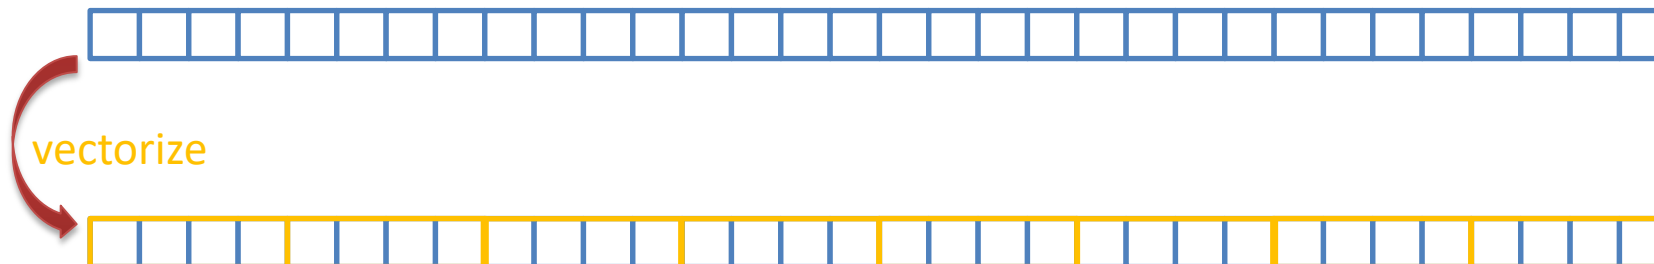
```
#pragma omp simd [clause[[, clause],...]  
for-loops
```

■ Syntax (Fortran)

```
!$omp simd [clause[[, clause],...]  
do-loops  
[!$omp end simd]
```

Example

```
float sprod(float *a, float *b, int n) {  
    float sum = 0.0f;  
    #pragma omp simd reduction(+:sum)  
    for (int k=0; k<n; k++)  
        sum += a[k] * b[k];  
    return sum;  
}
```



Data Sharing Clauses

- `private(var-list) :`

Uninitialized vectors for variables in *var-list*



- `firstprivate(var-list) :`

Initialized vectors for variables in *var-list*



- `reduction(op:var-list) :`

Create private variables for *var-list* and apply reduction operator *op* at the end of the construct



SIMD Loop Clauses

■ `safelen (length)`

- Maximum number of iterations that can run concurrently without breaking a dependence
- In practice, maximum vector length

■ `linear (list[:linear-step])`

- The variable's value is in relationship with the iteration number
 - $x_i = x_{\text{orig}} + i * \text{linear-step}$

■ `aligned (list[:alignment])`

- Specifies that the list items have a given alignment
- Default is alignment for the architecture

■ `collapse (n)`

SIMD Worksharing Construct

■ Parallelize and vectorize a loop nest

- Distribute a loop's iteration space across a thread team
- Subdivide loop chunks to fit a SIMD vector register

■ Syntax (C/C++)

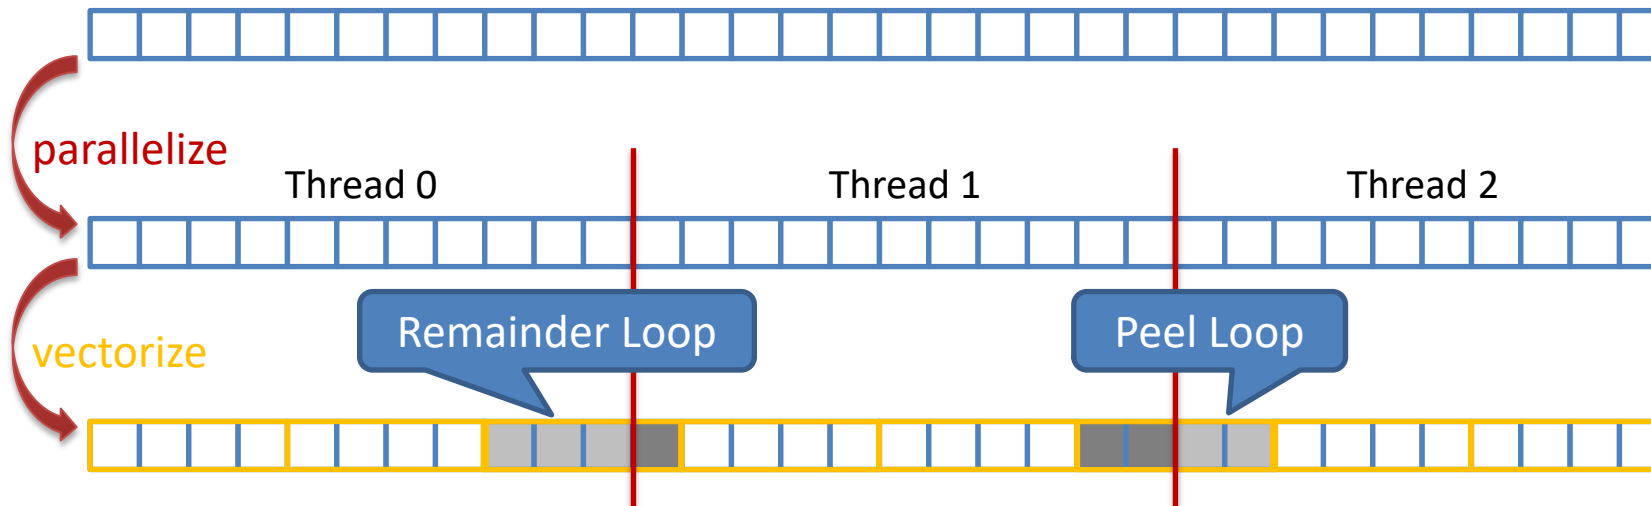
```
#pragma omp for simd [clause[[, clause],...]  
for-loops
```

■ Syntax (Fortran)

```
!$omp do simd [clause[[, clause],...]  
do-loops  
[!$omp end do simd [nowait]]
```

Example

```
float sprod(float *a, float *b, int n) {  
    float sum = 0.0f;  
    #pragma omp for simd reduction(+:sum)  
    for (int k=0; k<n; k++)  
        sum += a[k] * b[k];  
    return sum;  
}
```



Be Careful What You Wish For...

```
float sprod(float *a, float *b, int n) {  
    float sum = 0.0f;  
    #pragma omp for simd reduction(+:sum) \  
                                schedule(static, 5)  
    for (int k=0; k<n; k++)  
        sum += a[k] * b[k];  
    return sum;  
}
```

- You should choose chunk sizes that are multiples of the SIMD length
 - Remainder loops are not triggered
 - Likely better performance
- In the above example ...
 - and AVX2, the code will only execute the remainder loop!
 - and SSE, the code will have one iteration in the SIMD loop plus one in the remainder loop!

OpenMP 4.5 Simplifies SIMD Chunks

```
float sprod(float *a, float *b, int n) {  
    float sum = 0.0f;  
    #pragma omp for simd reduction(+:sum) \  
                                schedule(simd: static, 5)  
    for (int k=0; k<n; k++)  
        sum += a[k] * b[k];  
    return sum;  
}
```

- Chooses chunk sizes that are multiples of the SIMD length
 - First and last chunk may be slightly different to fix alignment and to handle loops that are not exact multiples of SIMD width
 - Remainder loops are not triggered
 - Likely better performance

SIMD Function Vectorization

```
float min(float a, float b) {  
    return a < b ? a : b;  
}  
  
float distsq(float x, float y) {  
    return (x - y) * (x - y);  
}  
  
void example() {  
    #pragma omp parallel for simd  
    for (i=0; i<N; i++) {  
        d[i] = min(distsq(a[i], b[i]), c[i]);  
    }  
}
```

SIMD Function Vectorization

- Declare one or more functions to be compiled for calls from a SIMD-parallel loop

- Syntax (C/C++):

```
#pragma omp declare simd [clause[[, clause],...]  
[#pragma omp declare simd [clause[[, clause],...]]  
[...]  
function-definition-or-declaration
```

- Syntax (Fortran):

```
!$omp declare simd (proc-name-list)
```

SIMD Function Vectorization

```
#pragma omp declare simd
```

```
float min(float a, float b) {  
    return a < b ? a : b;  
}
```

```
_ZGVZN16vv_min(%xmm0, %xmm1):  
    vminps %xmm1, %xmm0, %xmm0  
    ret
```

```
#pragma omp declare simd
```

```
float distsq(float x, float y)  
    return (x - y) * (x - y);  
}
```

```
_ZGVZN16vv_distsq(%xmm0, %xmm1):  
    vsubps %xmm0, %xmm1, %xmm2  
    vmulps %xmm2, %xmm2, %xmm0  
    ret
```

```
void example() {
```

```
#pragma omp parallel for simd
```

```
    for (i=0; i<N; i++) {  
        d[i] = min(distsq(a[i], b[i]), c[i]);  
    } }
```

```
vmovups (%r14,%r12,4), %xmm0  
vmovups (%r13,%r12,4), %xmm1  
call _ZGVZN16vv_distsq  
vmovups (%rbx,%r12,4), %xmm1  
call _ZGVZN16vv_min
```


SIMD Function Vectorization

- `simdlen (length)`

- generate function to support a given vector length

- `uniform (argument-list)`

- argument has a constant value between the iterations of a given loop

- `inbranch`

- function always called from inside an if statement

- `notinbranch`

- function never called from inside an if statement

- `linear (argument-list[:linear-step])`


- `aligned (argument-list[:alignment])`

inbranch & notinbranch

```
#pragma omp declare simd inbranch
```


```
float do_stuff(float x) {  
    /* do something */  
    return x * 2.0;  
}
```

```
vec8 do_stuff_v(vec8 x, mask m) {  
    /* do something */  
    vmulpd x{m}, 2.0, tmp  
    return tmp;  
}
```

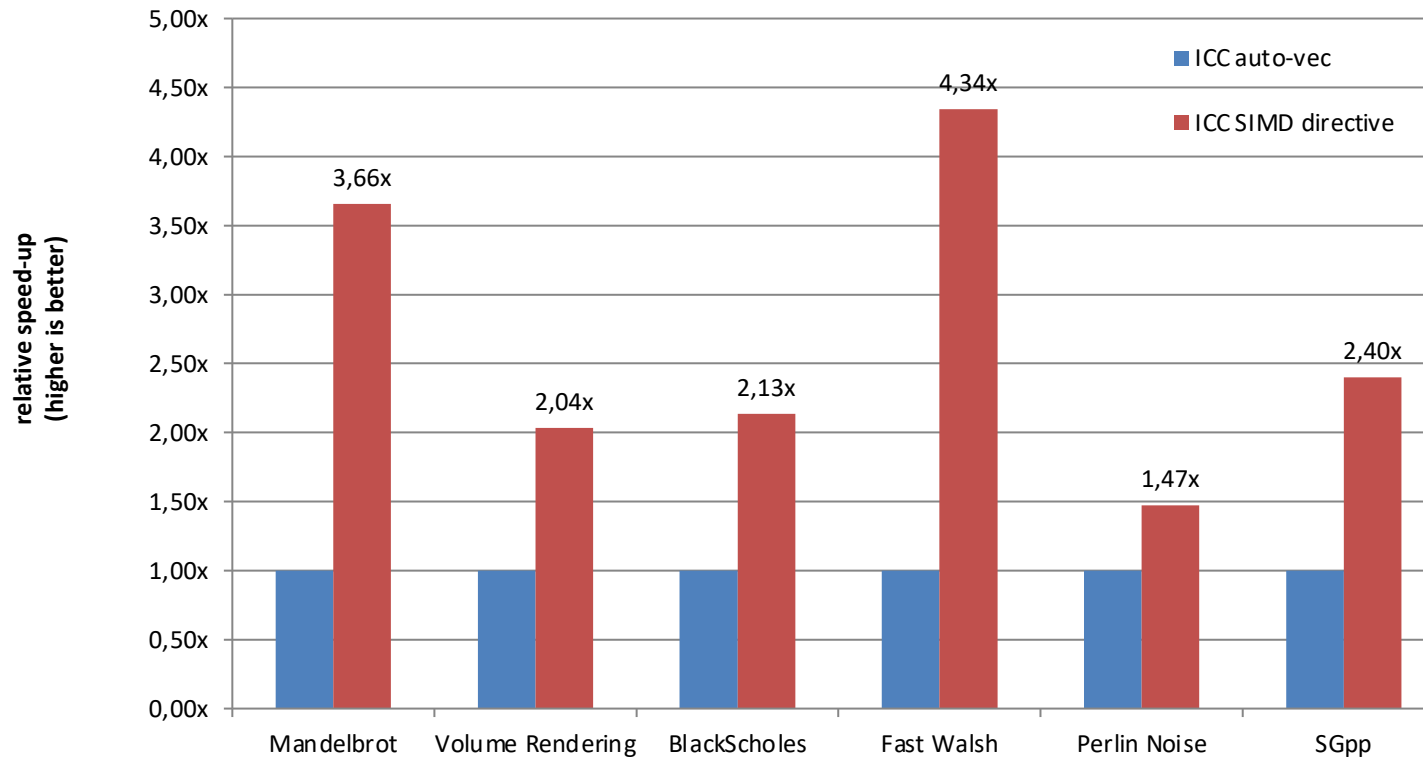


```
void example() {  
    #pragma omp simd  
    for (int i = 0; i < N; i++)  
        if (a[i] < 0.0)  
            b[i] = do_stuff(a[i]);  
}
```

```
for (int i = 0; i < N; i+=8) {  
    vcmp_lt &a[i], 0.0, mask  
    b[i] = do_stuff_v(&a[i], mask);  
}
```



SIMD Constructs & Performance



M.Klemm, A.Duran, X.Tian, H.Saito, D.Caballero, and X.Martorell. Extending OpenMP with Vector Constructs for Modern Multicore SIMD Architectures. In Proc. of the Intl. Workshop on OpenMP, pages 59-72, Rome, Italy, June 2012. LNCS 7312.

OpenMP Offload Programming

Introduction to OpenMP Offload Features

Running Example for this Presentation: saxpy

```
void saxpy() {  
    float a, x[SZ], y[SZ];  
    // left out initialization  
    double t = 0.0;  
    double tb, te;  
    tb = omp_get_wtime();  
    #pragma omp parallel for firstprivate(a)  
    for (int i = 0; i < SZ; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
    te = omp_get_wtime();  
    t = te - tb;  
    printf("Time of kernel: %lf\n", t);  
}
```

Timing code (not needed, just to have a bit more code to show 😊)

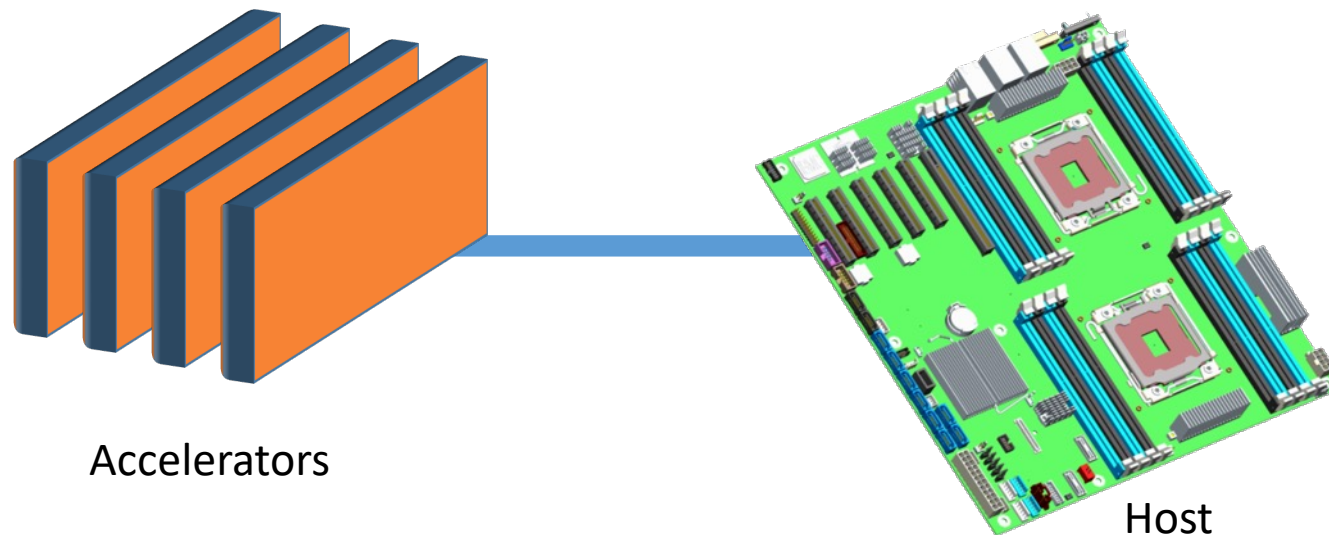
This is the code we want to execute on a target device (i.e., GPU)

Timing code (not needed, just to have a bit more code to show 😊)

Don't do this at home!
Use a BLAS library for this!

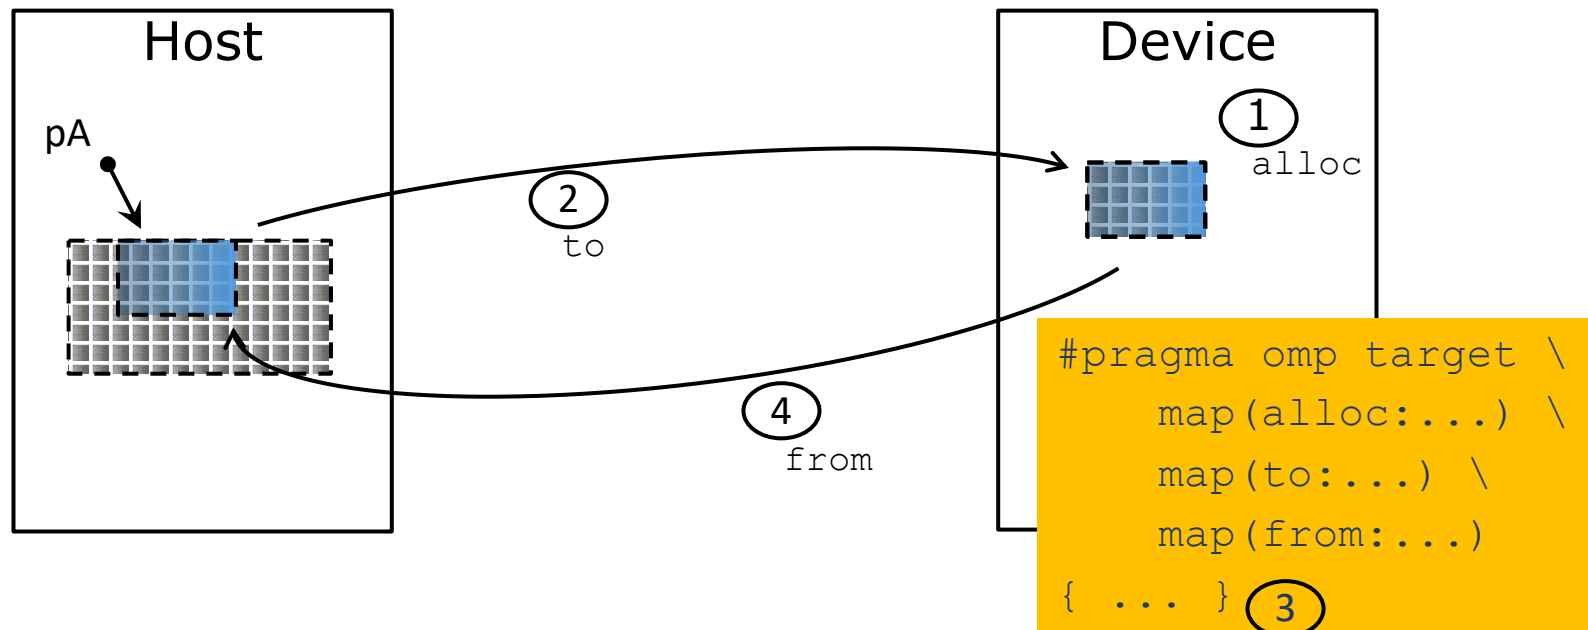
Device Model

- As of version 4.0 the OpenMP API supports accelerators/coprocessors
- Device model:
 - One host for “traditional” multi-threading
 - Multiple accelerators/coprocessors of the same kind for offloading



Execution Model

- Offload region and data environment is lexically scoped
 - Data environment is destroyed at closing curly brace
 - Allocated buffers/data are automatically released



OpenMP for Devices - Constructs

- Transfer control and data from the host to the device

- Syntax (C/C++)

```
#pragma omp target [clause[[, clause],...]  
structured-block
```

- Syntax (Fortran)

```
!$omp target [clause[[, clause],...]  
structured-block  
!$omp end target
```

- Clauses

```
device(scalar-integer-expression)  
map([{alloc | to | from | tofrom}]:) list)  
if(scalar-expr)
```

Example: saxpy

```
void saxpy() {  
    float a, x[SZ], y[SZ];  
    double t = 0.0;  
    double tb, te;  
    tb = omp_get_wtime();  
    #pragma omp target "map(tofrom:y[0:SZ])"  
    for (int i = 0; i < SZ; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
    te = omp_get_wtime();  
    t = te - tb;  
    printf("Time of kernel: %lf\n", t);  
}
```

The compiler identifies variables that are used in the target region.

All accessed arrays are copied from host to device and back

a
x[0:SZ]
y[0:SZ]

target

Presence check: only transfer if not yet allocated on the device.

x[0:SZ]
y[0:SZ]

Copying x back is not necessary: it was not changed.

clang -fopenmp --offload-arch=gfx90a ...

Example: saxpy

```
subroutine saxpy(a, x, y, n)
  use iso_fortran_env
  integer :: n, i
  real(kind=real32) :: a
  real(kind=real32), dimension(n) :: x
  real(kind=real32), dimension(n) :: y
```

```
!$omp target "map(tofrom:y(1:n))"
```

```
do i=1,n
```

```
  y(i) = a * x(i) + y(i)
```

```
end do
```

```
!$omp end target
```

```
end subroutine
```

The compiler identifies variables that are used in the target region.

All accessed arrays are copied from host to device and back

a
x(1:n)
y(1:n)

Presence check: only transfer if not yet allocated on the device.

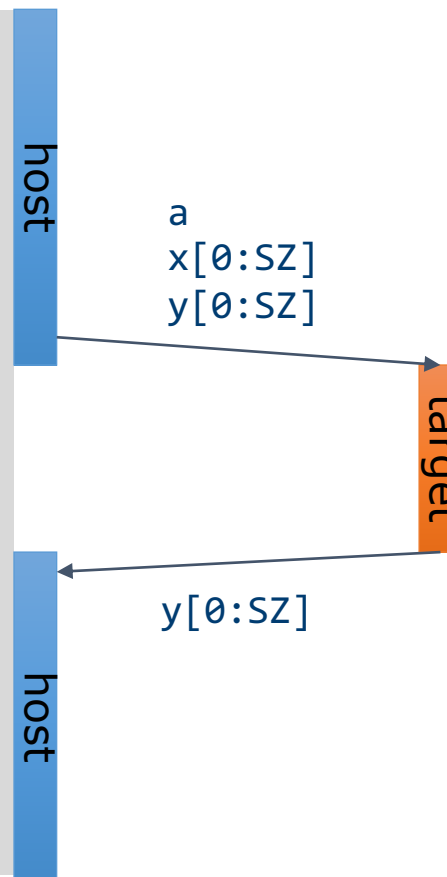
x(1:n)
y(1:n)

Copying x back is not necessary: it was not changed.

flang -fopenmp --offload-arch=gfx90a ...

Example: saxpy

```
void saxpy() {  
    double a, x[SZ], y[SZ];  
    double t = 0.0;  
    double tb, te;  
    tb = omp_get_wtime();  
    #pragma omp target map(to:x[0:SZ]) \  
                        map(tofrom:y[0:SZ])  
    for (int i = 0; i < SZ; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
    te = omp_get_wtime();  
    t = te - tb;  
    printf("Time of kernel: %lf\n", t);  
}
```



clang -fopenmp --offload-arch=gfx90a ...

Example: saxpy

```
void saxpy(float a, float* x, float* y,
           int sz) {
    double t = 0.0;
    double tb, te;
    tb = omp_get_wtime();
    #pragma omp target map(to:x[0:sz]) \
                      map(tofrom:y[0:sz])
    for (int i = 0; i < sz; i++) {
        y[i] = a * x[i] + y[i];
    }
    te = omp_get_wtime();
    t = te - tb;
    printf("Time of kernel: %lf\n", t);
}
```

The compiler cannot determine the size of memory behind the pointer.

host

a
x[0:sz]
y[0:sz]

target

host

y[0:sz]

Programmers have to help the compiler with the size of the data transfer needed.

```
clang -fopenmp --offload-arch=gfx90a ...
```

Creating Parallelism on the Target Device

- The `target` construct transfers the control flow to the target device
 - Transfer of control is sequential and synchronous
 - This is intentional!
- OpenMP separates offload and parallelism
 - Programmers need to explicitly create parallel regions on the target device
 - In theory, this can be combined with any OpenMP construct
 - In practice, there is only a useful subset of OpenMP features for a target device such as a GPU, e.g., no I/O, limited use of base language features.

Example: saxpy

```
void saxpy(float a, float* x, float* y,  
          int sz) {  
    #pragma omp target map(to:x[0:sz]) \  
                        map(tofrom:y[0:sz])  
    #pragma omp parallel for simd  
    for (int i = 0; i < sz; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
}
```

host
target
host

Create a team of threads to execute the loop in parallel using SIMD instructions.

GPUs are multi-level devices:
SIMD, threads, thread blocks

clang -fopenmp --offload-arch=gfx90a ...

teams Construct

- Support multi-level parallel devices

- Syntax (C/C++):

```
#pragma omp teams [clause[[,] clause],...]  
structured-block
```

- Syntax (Fortran):

```
!$omp teams [clause[[,] clause],...]  
structured-block
```

- Clauses

```
num_teams(integer-expression), thread_limit(integer-expression)  
default(shared | firstprivate | private none)  
private(list), firstprivate(list), shared(list), reduction(operator:list)
```


Multi-level Parallel saxpy

■ Manual code transformation

- Tile the loops into an outer loop and an inner loop
- Assign the outer loop to “teams” (OpenCL: work groups)
- Assign the inner loop to the “threads” (OpenCL: work items)

Multi-level Parallel saxpy

- For convenience, OpenMP defines composite constructs to implement the required code transformations

```
void saxpy(float a, float* x, float* y, int n) {  
    #pragma omp target teams distribute parallel for simd \  
        num_teams(num_blocks) map(to:x[0:sz]) map(tofrom:y[0:sz])  
    for (int i = 0; i < n; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
}
```

```
subroutine saxpy(a, x, y, n)  
    ! Declarations omitted  
!$omp omp target teams distribute parallel do simd &  
!$omp&        num_teams(num_blocks) map(to:x) map(tofrom:y)  
    do i=1,n  
        y(i) = a * x(i) + y(i)  
    end do  
!$omp end target teams distribute parallel do simd  
end subroutine
```

Optimize Data Transfers

■ Reduce the amount of time spent transferring data

- Use map clauses to enforce direction of data transfer.
- Use target data, target enter data, target exit data constructs to keep data environment on the target device.

```
void example() {  
    float tmp[N], a_in[N], b[N], c[N];  
    #pragma omp target data map(alloc:tmp[:N]) \  
        map(to:a[:N],b[:N]) \  
        map(tofrom:c[:N])  
  
    {  
        zeros(tmp, N);  
        compute_kernel_1(tmp, a, N); // uses target  
        saxpy(2.0f, tmp, b, N);  
        compute_kernel_2(tmp, b, N); // uses target  
        saxpy(2.0f, c, tmp, N);  
    }  
}
```

```
void zeros(float* a, int n) {  
    #pragma omp target teams distribute parallel for  
        for (int i = 0; i < n; i++)  
            a[i] = 0.0f;  
}
```

```
void saxpy(float a, float* y, float* x, int n) {  
    #pragma omp target teams distribute parallel for  
        for (int i = 0; i < n; i++)  
            y[i] = a * x[i] + y[i];  
}
```

target data Construct Syntax

- Create scoped data environment and transfer data from the host to the device and back

- Syntax (C/C++)

```
#pragma omp target data [clause[[, clause],...]  
structured-block
```

- Syntax (Fortran)

```
!$omp target data [clause[[, clause],...]  
structured-block  
!$omp end target data
```

- Clauses

```
device(scalar-integer-expression)  
map([{alloc | to | from | tofrom | release | delete}:] list)  
if(scalar-expr)
```

target update Construct Syntax

- Issue data transfers to or from existing data device environment

- Syntax (C/C++)

```
#pragma omp target update [clause[[, clause],...]
```

- Syntax (Fortran)

```
!$omp target update [clause[[, clause],...]
```

- Clauses

```
device(scalar-integer-expression)  
to(list)  
from(list)  
if(scalar-expr)
```

Example: target data and target update

```
#pragma omp target data device(0) map(alloc:tmp[:N]) map(to:input[:N]) map(from:res)
{
#pragma omp target device(0)
#pragma omp parallel for
    for (i=0; i<N; i++)
        tmp[i] = some_computation(input[i], i);

    update_input_array_on_the_host(input);

#pragma omp target update device(0) to(input[:N])

#pragma omp target device(0)
#pragma omp parallel for reduction(+:res)
    for (i=0; i<N; i++)
        res += final_computation(input[i], tmp[i], i)
}
```

host

target

host

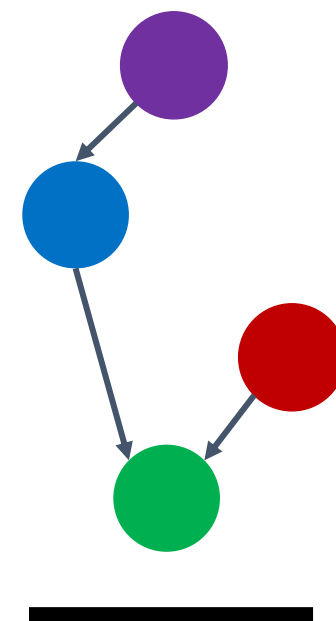
target

host

Asynchronous Offloads

- OpenMP target constructs are synchronous by default
 - The encountering host thread awaits the end of the target region before continuing
 - The `nowait` clause makes the target constructs asynchronous (in OpenMP speak: they become an OpenMP task)

<code>#pragma omp task</code>		<code>depend(out:a)</code>
<code>init_data(a);</code>		
<code>#pragma omp target map(to:a[:N]) map(from:x[:N])</code>	<code>nowait</code>	<code>depend(in:a) depend(out:x)</code>
<code>compute_1(a, x, N);</code>		
<code>#pragma omp target map(to:b[:N]) map(from:y[:N])</code>	<code>nowait</code>	<code>depend(out:y)</code>
<code>compute_2(b, y, N);</code>		
<code>#pragma omp target map(to:y[:N]) map(to:x[:N])</code>	<code>nowait</code>	<code>depend(in:x) depend(in:y)</code>
<code>compute_3(x, y, N);</code>		
<code>#pragma omp taskwait</code>		



Hybrid Programming

Hybrid Programming

- Hybrid programming here stands for the interaction of OpenMP with a lower-level programming model, e.g.
 - OpenCL
 - CUDA
 - HIP

- OpenMP supports these interactions
 - Calling low-level kernels from OpenMP application code
 - Calling OpenMP kernels from low-level application code

Example: Calling saxpy

```
void example() {  
    float a = 2.0;  
    float * x;  
    float * y;  
  
    // allocate the device memory  
    #pragma omp target data map(to:x[0:count]) map(tofrom:y[0:count])  
    {  
        compute_1(n, x);  
        compute_2(n, y);  
        saxpy(n, a, x, y)  
        compute_3(n, y);  
    }  
}
```

Let's assume that we want to implement the saxpy() function in a low-level language.

```
void saxpy(size_t n, float a,  
           float * x, float * y) {  
    #pragma omp target teams distribute \  
        parallel for simd  
    for (size_t i = 0; i < n; ++i) {  
        y[i] = a * x[i] + y[i];  
    }  
}
```

HIP Kernel for saxpy()

- Assume a HIP version of the SAXPY kernel:

```
__global__ void saxpy_kernel(size_t n, float a, float * x, float * y) {  
    size_t i = threadIdx.x + blockIdx.x * blockDim.x;  
    y[i] = a * x[i] + y[i];  
}
```

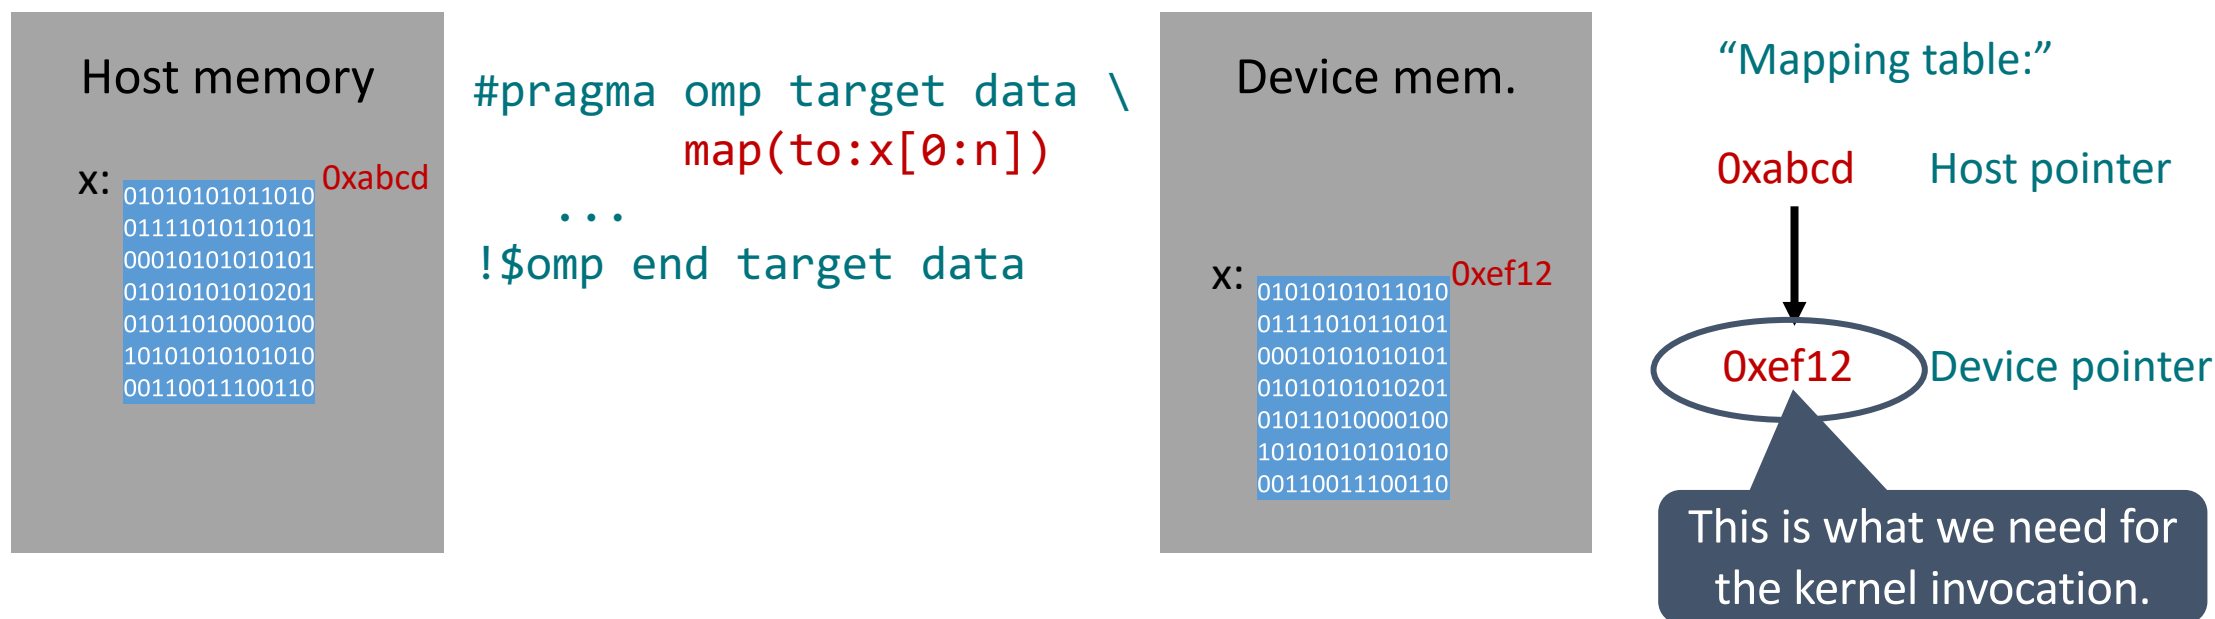
```
void saxpy_hip(size_t n, float a, float * x, float * y) {  
    assert(n % 256 == 0);  
    saxpy_kernel<<<n/256,256,0,NULL>>>(n, a, x, y);  
}
```

These are device pointers!

- We need a way to translate the host pointer that was mapped by OpenMP directives and retrieve the associated device pointer.

Pointer Translation /1

- When creating the device data environment, OpenMP creates a mapping between
 - the (virtual) memory pointer on the host and
 - the (virtual) memory pointer on the target device.
- This mapping is established through the data-mapping directives and their clauses.



Pointer Translation /2

- The target data construct defines the `use_device_addr` clause to perform pointer translation.
 - The OpenMP implementation searches for the host pointer in its internal mapping tables.
 - The associated device pointer is then returned.

```
type * x = 0xabcd;  
#pragma omp target data use_device_addr(x[0:0])  
{  
    example_func(x);    // x == 0xef12  
}
```

- Note: the pointer variable shadowed within the `target data` construct for the translation.

Putting it Together...

```
void example() {  
    float a = 2.0;  
    float * x = ...;    // assume: x = 0xabcd  
    float * y = ...;  
  
    // allocate the device memory  
    #pragma omp target data map(to:x[0:count]) map(tofrom:y[0:count])  
    {  
        compute_1(n, x); // mapping table: x:[0xabcd,0xef12], x = 0xabcd  
        compute_2(n, y);  
        #pragma omp target data use_device_addr(x[0:0],y[0:0])  
        {  
            saxpy_hip(n, a, x, y) // mapping table: x:[0xabcd,0xef12], x = 0xef12  
        }  
        compute_3(n, y);  
    }  
}
```

Summary

- OpenMP API is ready to use Intel discrete GPUs for offloading compute
 - Mature offload model w/ support for asynchronous offload/transfer
 - Tightly integrates with OpenMP multi-threading on the host
- More, advanced features (not covered here)
 - Memory management API
 - Interoperability with native data management
 - Interoperability with native streaming interfaces
 - Unified shared memory support



Visit www.openmp.org for more information

Programming OpenMP

GPU: expressing parallelism

Christian Terboven

Michael Klemm

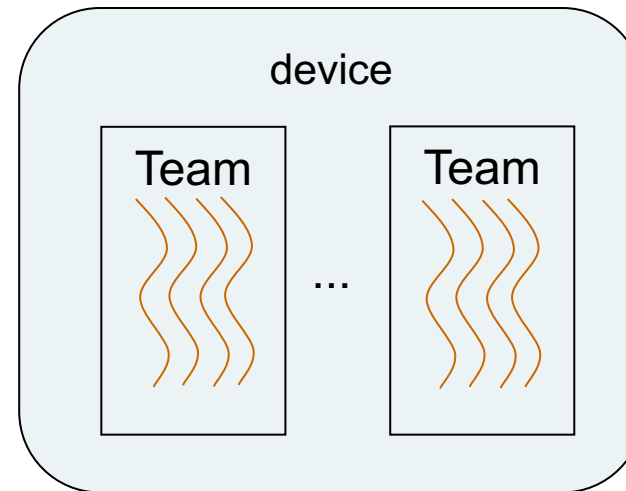


teams and distribute constructs

Many slides are taken from the lecture High-Performance Computing at RWTH Aachen University
Authors include: Sandra Wienke, Julian Miller

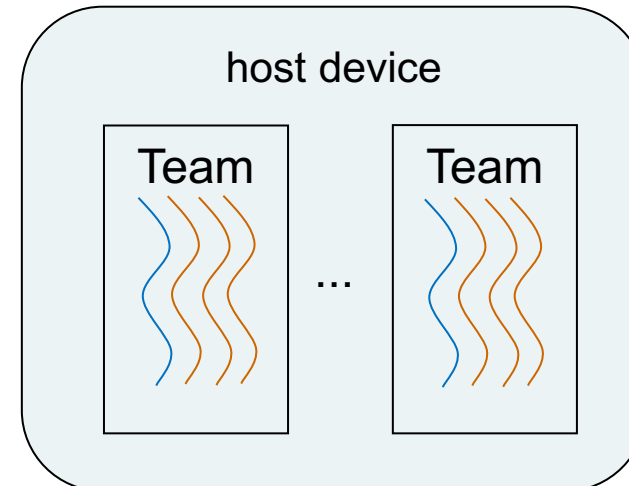
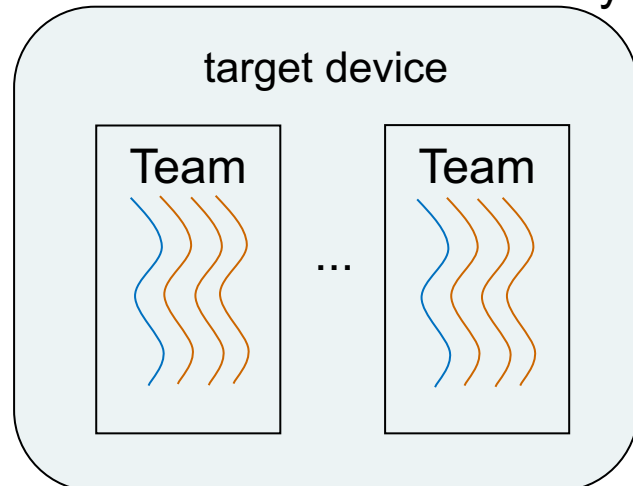
Terminology

- **League:**
the set of threads teams created by a **teams** construct
- **Contention group:**
threads of a team in a league and their descendant threads



The **teams** construct creates a *league* of thread teams

- The master thread of each team executes the **teams** region
- The number of teams is specified by the **num_teams** clause
- Each team executes with **thread_limit** threads
- Threads in different teams cannot synchronize with each other



5.0

Only special OpenMP constructs or routines can be strictly nested inside a **teams** construct:

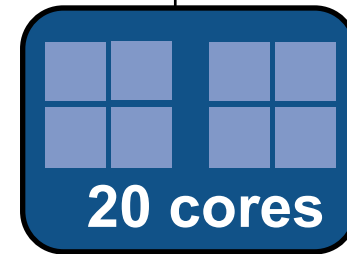
- **distribute** [**simd**], **distribute** [**parallel**] worksharing-loop [**SIMD**]
- **parallel** regions (**parallel for/do**, **parallel sections**)
- **omp_get_num_teams()** and **omp_get_team_num()**

distribute Construct

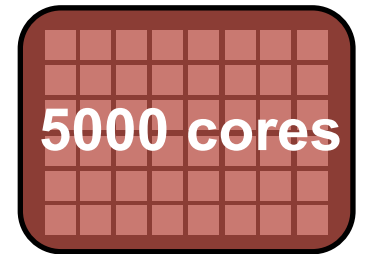
- work sharing among the teams regions
 - Distribute the iterations of the associated loops across the master threads of each team executing the region
- Strictly nested inside a teams region
- No implicit barrier at the end of the construct
- **dist_schedule**(*kind*[, *chunk_size*])
 - The scheduling kind must be **static**
 - Chunks are distributed in round-robin fashion of chunks with size *chunk_size*
 - If no chunk size specified, chunks are of (almost) equal size; each team receives at most one chunk

Example DAXPY: How to Port to GPU?

```
void daxpy(int n, double a, double *x, double *y) {  
    for (int i = 0; i < n; ++i)  
        y[i] = a * x[i] + y[i];  
}  
  
int main(int argc, const char* argv[]) {  
    static int n = 100000000; static double a = 2.0;  
    double *x = (double *) malloc(n * sizeof(double));  
    double *y = (double *) malloc(n * sizeof(double));  
  
    // Initialize x, y  
    for(int i = 0; i < n; ++i){  
        x[i] = 1.0;  
        y[i] = 2.0;  
    }  
    daxpy(n, a, x, y); // Invoke daxpy kernel  
    // Check if all values are 4.0  
  
    free(x); free(y);  
    return 0;  
}
```



CPU



GPU

How to port
DAXPY to a GPU?

Kernel Directives

- Offload kernel code
 - `target`: offload work
 - `teams, parallel`: create in parallelly running threads
 - `distribute, do, for, simd`: worksharing across parallel units
- Worksharing
 - `for`: offload work
 - `collapse`: collapse two or more nested loops to increase parallelism

```
clang -fopenmp -Xopenmp-target -fopenmp-targets=nvptx64-nvidia-cuda -march=sm_70  
--cuda-path=$CUDA_TOOLKIT_ROOT_DIR daxpy.c
```

- **clang** A recent clang compiler with OpenMP target support
- **-fopenmp** Enables general OpenMP support
- **-Xopenmp-target** Enables OpenMP target support
- **-fopenmp-targets=nvptx64-nvidia-cuda** Specifies the target architecture → here: NVIDIA GPUs
- **-march=sm_70** Optional. Specifies the target compute architecture
- **--cuda-path=\$CUDA_TOOLKIT_ROOT_DIR** Optional. Specifies the CUDA path

Example: DAXPY

```
void daxpy(int n, double a, double *x, double *y) {  
    #pragma omp target  
    for (int i = 0; i < n; ++i)  
        y[i] = a * x[i] + y[i];  
}
```

```
int main(int argc, const char* argv[]) {  
    static int n = 100000000; static double a = 2.0;  
    double *x = (double *) malloc(n * sizeof(double));  
    double *y = (double *) malloc(n * sizeof(double));
```

```
    // Initialize x, y  
    for(int i = 0; i < n; ++i){  
        x[i] = 1.0;  
        y[i] = 2.0;  
    }  
    daxpy(n, a, x, y); // Invoke daxpy kernel  
    // Check if all values are 4.0
```

```
    free(x); free(y);  
    return 0;  
}
```

Output:

```
$ $CC $FLAGS_OFFLOAD_OPENMP daxpy.c  
$ a.out
```

Libomptarget fatal error 1: failure of target
construct while offloading is mandatory

Example DAXPY: Debugging

- No compiler error but cryptic runtime error
- NVIDIA Profiler

```
$ nvprof daxpy.exe
==40419== NVPROF is profiling process 40419, command: daxpy.exe
==40419== Profiling application: daxpy.exe
==40419== Profiling result:
No kernels were profiled.

==40419== API calls:
No API activities were profiled.
```

- Cuda-memcheck

```
$ cuda-memcheck daxpy.exe
===== CUDA-MEMCHECK
===== Invalid __global__ read of size 8
===== at 0x00000d10 in __omp_offloading_4b_f850d140_daxpy_l3
===== by thread (32,0,0) in block (0,0,0)
===== Address 0x00000000 is out of bounds
```


Example DAXPY: Data Management

```
void daxpy(int n, double a, double *x, double *y) {  
    #pragma omp target map(tofrom:y[0:n]) map(to:a,x[0:n])  
    for (int i = 0; i < n; ++i)  
        y[i] = a * x[i] + y[i];  
}  
  
int main(int argc, const char* argv[]) {  
    static int n = 100000000; static double a = 2.0;  
    double *x = (double *) malloc(n * sizeof(double));  
    double *y = (double *) malloc(n * sizeof(double));  
  
    // Initialize x, y  
    for(int i = 0; i < n; ++i){  
        x[i] = 1.0;  
        y[i] = 2.0;  
    }  
    daxpy(n, a, x, y); // Invoke daxpy kernel  
    // Check if all values are 4.0  
  
    free(x); free(y);  
    return 0;  
}
```

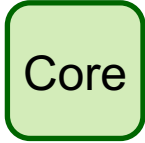
For comparison:
~0.12s on a
single CPU core

Output:
\$ \$CC \$FLAGS_OFFLOAD_OPENMP daxpy.c
\$ a.out
Max error: 0.00000
Total runtime: 102.50s

Mapping to Hardware

Thread




Core


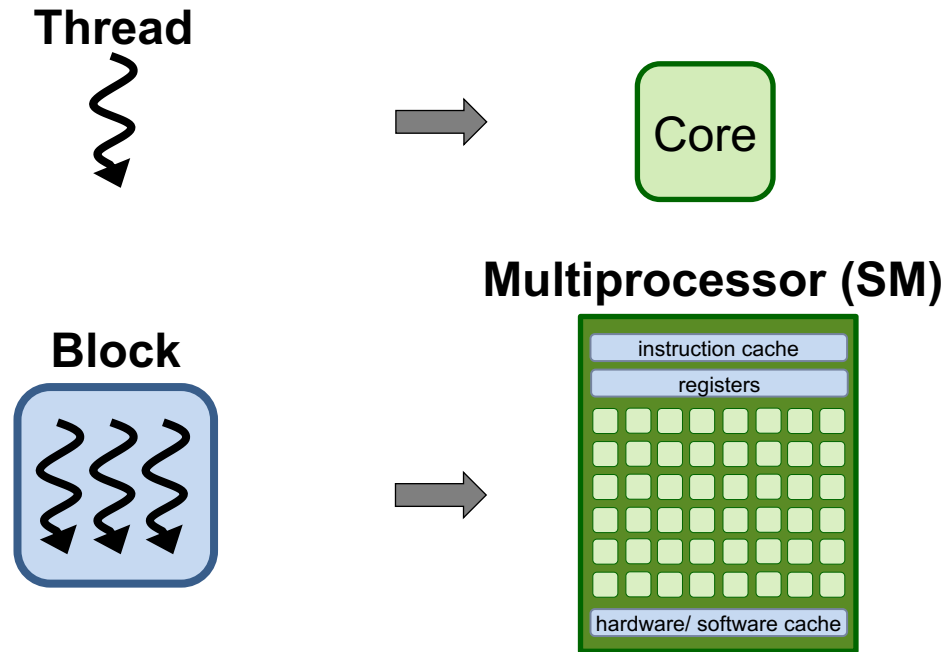
- Each thread is executed by a core

Example DAXPY: Thread Parallelism

```
void daxpy(int n, double a, double *x, double *y) {  
    #pragma omp target parallel for map(tofrom:y[0:n]) map(to:a,x[0:n])  
    for (int i = 0; i < n; ++i)  
        y[i] = a * x[i] + y[i];  
}  
  
int main(int argc, const char* argv[]) {  
    static int n = 100000000; static double a = 2.0;  
    double *x = (double *) malloc(n * sizeof(double));  
    double *y = (double *) malloc(n * sizeof(double));  
  
    // Initialize x, y  
    for(int i = 0; i < n; ++i){  
        x[i] = 1.0;  
        y[i] = 2.0;  
    }  
    daxpy(n, a, x, y); // Invoke daxpy kernel  
    // Check if all values are 4.0  
  
    free(x); free(y);  
    return 0;  
}
```

Output:
\$ \$CC \$FLAGS_OFFLOAD_OPENMP daxpy.c
\$ a.out
Max error: 0.00000
Total runtime: 9.65s

Mapping to Hardware



- Each thread is executed by a core
- Each block is executed on a SM
- Several concurrent blocks can reside on a SM depending on shared resources

Example DAXPY: Thread Parallelism

```
void daxpy(int n, double a, double *x, double *y) {  
    #pragma omp target teams distribute parallel for map(tofrom:y[0:n]) map(to:a,x[0:n])  
    for (int i = 0; i < n; ++i)  
        y[i] = a * x[i] + y[i];  
}  
  
int main(int argc, const char* argv[]) {  
    static int n = 100000000; static double a = 2.0;  
    double *x = (double *) malloc(n * sizeof(double));  
    double *y = (double *) malloc(n * sizeof(double));  
  
    // Initialize x, y  
    for(int i = 0; i < n; ++i){  
        x[i] = 1.0;  
        y[i] = 2.0;  
    }  
    daxpy(n, a, x, y); // Invoke daxpy kernel  
    // Check if all values are 4.0  
  
    free(x); free(y);  
    return 0;  
}
```

Output:

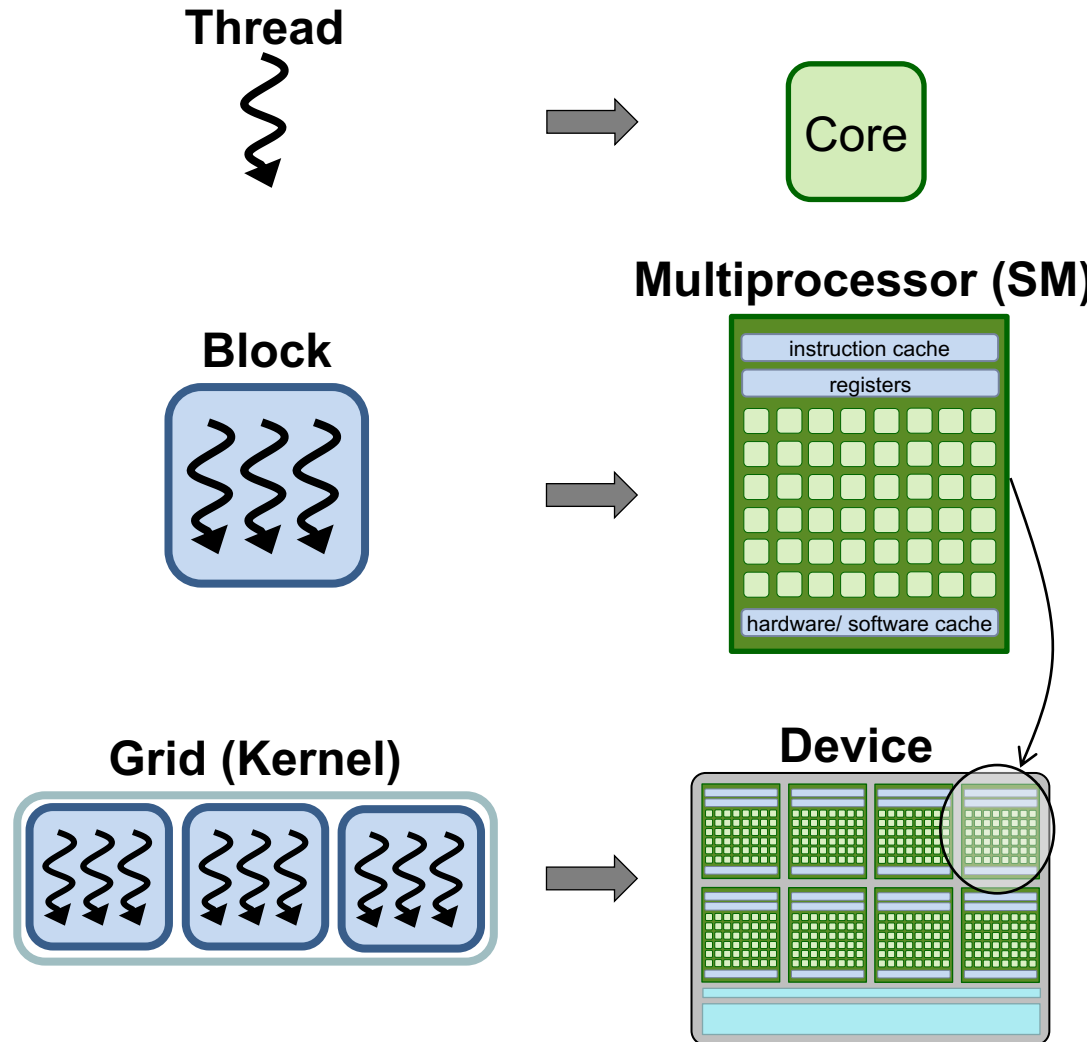
\$ \$CC \$FLAGS_OFFLOAD_OPENMP daxpy.c

\$ a.out

Max error: 0.00000

Total runtime: 0.80s

Mapping to Hardware



- Each thread is executed by a core
- Each block is executed on a SM
- Several concurrent blocks can reside on a SM depending on shared resources
- Each kernel is executed on a device

teams Construct

- **Syntax (C/C++):**
`#pragma omp teams [clause[[,] clause]...]
 structured-block`
- **Syntax (Fortran):**
`!$omp teams [clause[[,] clause]...]
 structured-block`
- **Clauses**
`num_teams(integer-expression)
thread_limit(integer-expression)
default(shared | none) OR
default(shared|private|firstprivate|none)
private(list)
firstprivate(list)
shared(list)
reduction([default,]reduction-identifier : list)
allocate([allocator:]list)`

5.0

distribute Construct

- Syntax (C/C++):
`#pragma omp distribute [clause[,] clause...]
for-loops`
- Syntax (Fortran):
`!$omp distribute [clause[,] clause...]
do-loops`
- Clauses
`private(list)`
`firstprivate(list)`
`lastprivate(list)`
`collapse(n)`
`dist_schedule(kind[, chunk_size])`
`allocate([allocator:]list)`

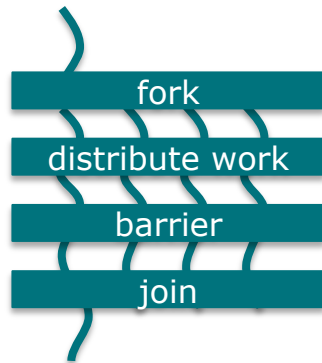
5.0

OpenMP Parallel Loops

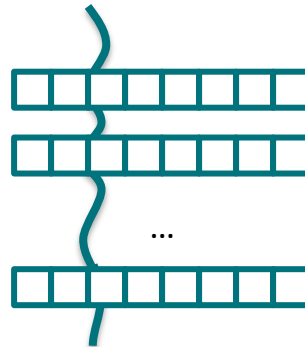
loop Construct

- Existing loop constructs are tightly bound to execution model:

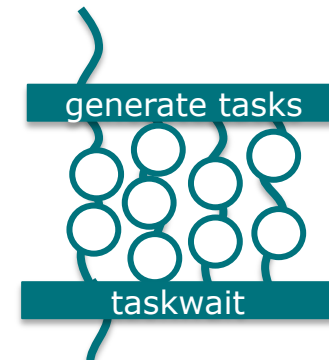
```
#pragma omp parallel for  
for (i=0; i<N;++i) {...}
```



```
#pragma omp simd  
for (i=0; i<N;++i) {...}
```



```
#pragma omp taskloop  
for (i=0; i<N;++i) {...}
```



- The `loop` construct is meant to tell OpenMP about truly parallel semantics of a loop.

OpenMP Fully Parallel Loops

```
int main(int argc, const char* argv[]) {  
    float *x = (float*) malloc(n * sizeof(float));  
    float *y = (float*) malloc(n * sizeof(float));  
    // Define scalars n, a, b & initialize x, y  
  
    #pragma omp parallel  
    #pragma omp loop  
        for (int i = 0; i < n; ++i){  
            y[i] = a*x[i] + y[i];  
        }  
    }  
}
```

Loop Constructs, Syntax

■ Syntax (C/C++)

```
#pragma omp loop [clause[[, clause],...]  
for-loops
```

■ Syntax (Fortran)

```
!$omp loop [clause[[, clause],...]  
do-loops  
[!$omp end loop]
```

Loop Constructs, Clauses

■ `bind(binding)`

- Binding region the loop construct should bind to
- One of: `teams`, `parallel`, `thread`

■ `order(concurrent)`

- Tell the OpenMP compiler that the loop can be executed in any order.
- Default!

■ `collapse(n)`

■ `private(list)`

■ `lastprivate(list)`

■ `reduction(reduction-id: list)`

Extensions to Existing Constructs

- Existing loop constructs have been extended to also have truly parallel semantics.

- C/C++ Worksharing:

```
#pragma omp [for|simd] order(concurrent) \  
[clause[[, clause],...]  
  
for-loops
```

- Fortran Worksharing:

```
!$omp [do|simd] order(concurrent) &  
[clause[[, clause],...]  
  
do-loops  
[!$omp end [do|simd]]
```


DOACROSS Loops

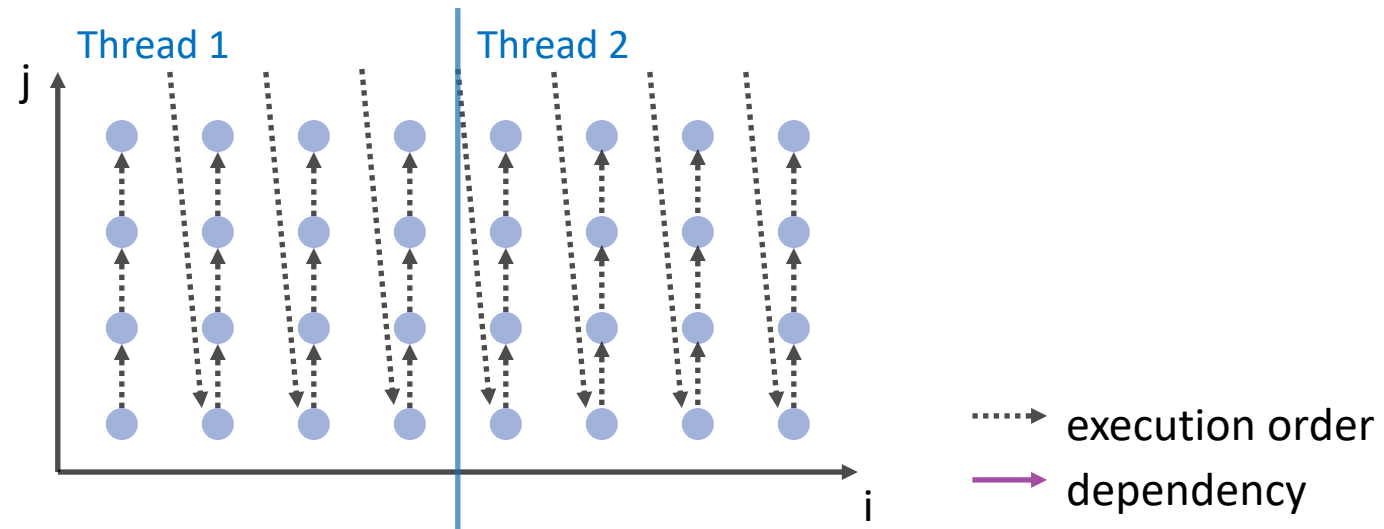
DOACROSS Loops

- “DOACROSS” loops are loops with special loop schedules
 - Restricted form of loop-carried dependencies
 - Require fine-grained synchronization protocol for parallelism
- Loop-carried dependency:
 - Loop iterations depend on each other
 - Source of dependency must be scheduled before sink of the dependency
- DOACROSS loop:
 - Data dependency is an invariant for the execution of the whole loop nest

Parallelizable Loops

- A parallel loop cannot not have any loop-carried dependencies (simplified just a little bit!)

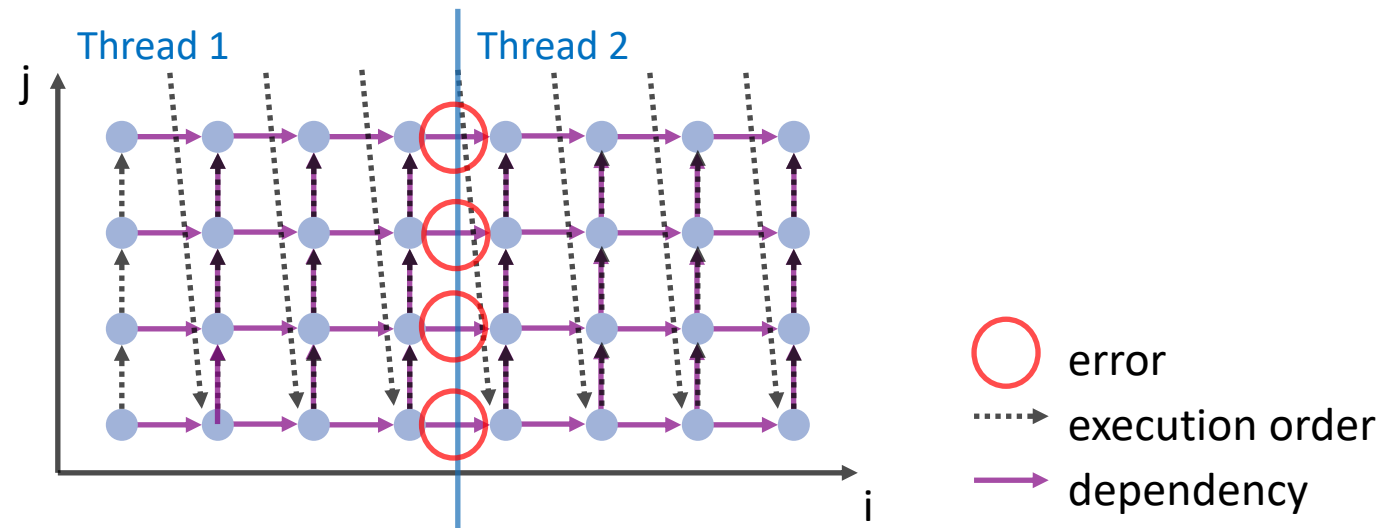
```
for (int i = 1; i < N; ++i) {  
    for (int j = 1; j < M; ++j) {  
        b[i][j] = f(b[i][j],  
                    b[i][j], a[i][j]);  
    }  
}
```



Non-parallelizable Loops

- If there is a loop-carried dependency, a loop cannot be parallelized anymore (“easily” that is)

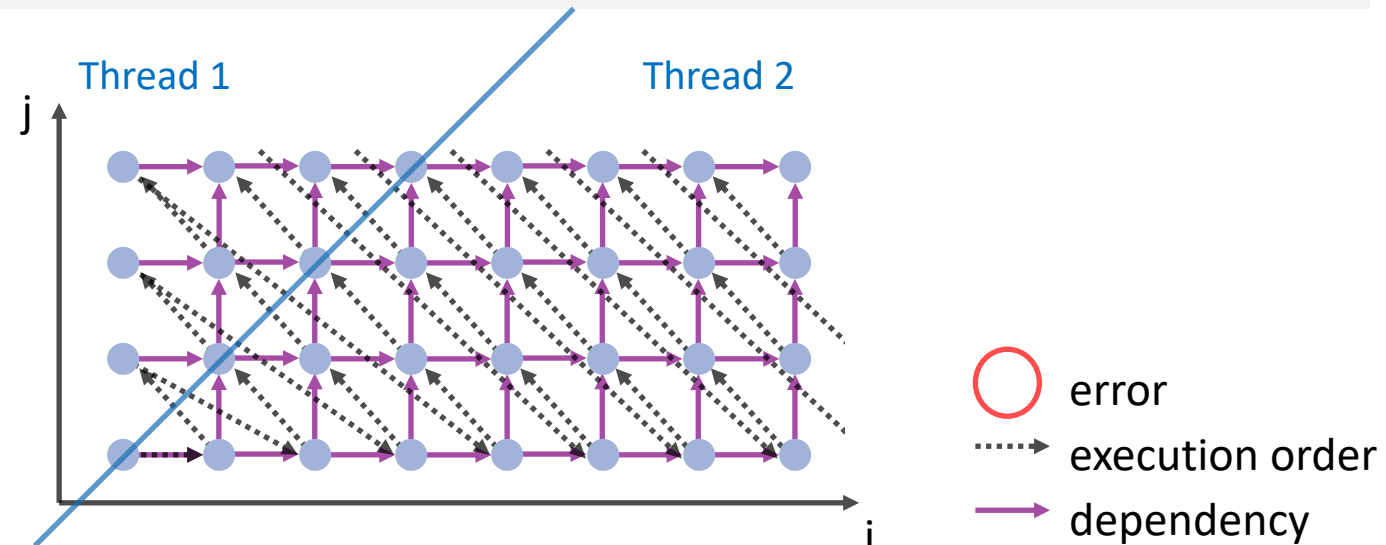
```
for (int i = 1; i < N; ++i) {  
    for (int j = 1; j < M; ++j) {  
        b[i][j] = f(b[i-1][j],  
                   b[i][j-1], a[i][j]);  
    }  
}
```



Wavefront-Parallel Loops

- If the data dependency is invariant, then skewing the loop helps remove the data dependency

```
for (int i = 1; i < N; ++i) {  
    for (int j = i+1; j < i+N; ++j) {  
        b[i][j-i] = f(b[i-1][j-i],  
                      b[i][j-i-1], a[i][j]);  
    }  
}
```



DOACROSS Loops with OpenMP

Deprecated
in v5.2

- OpenMP 4.5 extends the notion of the ordered construct to describe loop-carried dependencies

- Syntax (C/C++):

```
#pragma omp for ordered(d) [clause[[, clause],...]  
for-loops
```

and

```
#pragma omp ordered [clause[[, clause],...]
```

where *clause* is one of the following:

`depend(source)`

`depend(sink:vector)`

- Syntax (Fortran):

```
!$omp do ordered(d) [clause[[, clause],...]  
do-loops
```

```
!$omp ordered [clause[[, clause],...]
```

Example



- The ordered clause tells the compiler about loop-carried dependencies and their distances

```
#pragma omp parallel for ordered(2)
for (int i = 1; i < N; ++i) {
    for (int j = 1; j < M; ++j) {
        #pragma omp ordered depend(sink:i-1,j) depend(sink:i,j-1)
        b[i][j] = f(b[i-1][j],
                   b[i][j-1], a[i][j]);
    }
    #pragma omp ordered depend(source)
}
```

Example: 3D Gauss-Seidel

Deprecated
in v5.2

```
#pragma omp for ordered(2) private(j,k)
for (i = 1; i < N-1; ++i) {
    for (j = 1; j < N-1; ++j) {
        #pragma omp ordered depend(sink: i-1,j-1) depend(sink: i-1,j) \
            depend(sink: i-1,j+1) depend(sink: i,j-1)
        for (k = 1; k < N-1; ++k) {
            double tmp1 = (p[i-1][j-1][k-1] + p[i-1][j-1][k] + p[i-1][j-1][k+1]
                + p[i-1][j][k-1] + p[i-1][j][k] + p[i-1][j][k+1]
                + p[i-1][j+1][k-1] + p[i-1][j+1][k] + p[i-1][j+1][k+1]);
            double tmp2 = (p[i][j-1][k-1] + p[i][j-1][k] + p[i][j-1][k+1]
                + p[i][j][k-1] + p[i][j][k] + p[i][j][k+1]
                + p[i][j+1][k-1] + p[i][j+1][k] + p[i][j+1][k+1]);
            double tmp3 = (p[i+1][j-1][k-1] + p[i+1][j-1][k] + p[i+1][j-1][k+1]
                + p[i+1][j][k-1] + p[i+1][j][k] + p[i+1][j][k+1]
                + p[i+1][j+1][k-1] + p[i+1][j+1][k] + p[i+1][j+1][k+1]);
            p[i][j][k] = (tmp1 + tmp2 + tmp3) / 27.0;
        }
    }
    #pragma omp ordered depend(source)
}
}
```


DOACROSS Loops with OpenMP

- OpenMP 4.5 extends the notion of the ordered construct to describe loop-carried dependencies

- Syntax (C/C++):

```
#pragma omp for ordered [clause[[, clause],...]  
for-loops
```

and

```
#pragma omp ordered [clause[[, clause],...]
```

where *clause* is one of the following:

```
doacross(source:vector), vector can be omp_cur_iteration  
doacross(sink:vector)
```

- Syntax (Fortran):

```
!$omp do ordered [clause[[, clause],...]  
do-loops
```

```
!$omp ordered [clause[[, clause],...]
```

Example

- The ordered clause tells the compiler about loop-carried dependencies and their distances

```
#pragma omp parallel for ordered
for (int i = 1; i < N; ++i) {
    for (int j = 1; j < M; ++j) {
        #pragma omp ordered doacross(sink:i-1,j) doacross(sink:i,j-1)
        b[i][j] = f(b[i-1][j],
                    b[i][j-1], a[i][j]);
    }
    #pragma omp ordered doacross(source:omp_cur_iteration)
}
```

Example: 3D Gauss-Seidel

```
#pragma omp for ordered private(j,k)
for (i = 1; i < N-1; ++i) {
    for (j = 1; j < N-1; ++j) {
        #pragma omp ordered doacross(sink: i-1,j-1) doacross(sink: i-1,j) \
            doacross(sink: i-1,j+1) doacross(sink: i,j-1)
        for (k = 1; k < N-1; ++k) {
            double tmp1 = (p[i-1][j-1][k-1] + p[i-1][j-1][k] + p[i-1][j-1][k+1]
                          + p[i-1][j][k-1] + p[i-1][j][k] + p[i-1][j][k+1]
                          + p[i-1][j+1][k-1] + p[i-1][j+1][k] + p[i-1][j+1][k+1]);
            double tmp2 = (p[i][j-1][k-1] + p[i][j-1][k] + p[i][j-1][k+1]
                          + p[i][j][k-1] + p[i][j][k] + p[i][j][k+1]
                          + p[i][j+1][k-1] + p[i][j+1][k] + p[i][j+1][k+1]);
            double tmp3 = (p[i+1][j-1][k-1] + p[i+1][j-1][k] + p[i+1][j-1][k+1]
                          + p[i+1][j][k-1] + p[i+1][j][k] + p[i+1][j][k+1]
                          + p[i+1][j+1][k-1] + p[i+1][j+1][k] + p[i+1][j+1][k+1]);
            p[i][j][k] = (tmp1 + tmp2 + tmp3) / 27.0;
        }
    }
    #pragma omp ordered depend(source:omp_cur_iteration)
}
}
```

OpenMP Meta-Programming

The metadirective Directive

- Construct OpenMP directives for different OpenMP contexts
- Limited form of meta-programming for OpenMP directives and clauses

```
#pragma omp target map(to:v1,v2) map(from:v3)
#pragma omp metadirective \
    when( device={arch(nvptx)}: teams loop ) \
    default( parallel loop )
for (i = lb; i < ub; i++)
    v3[i] = v1[i] * v2[i];
```

```
!$omp begin metadirective &
    when( implementation={unified_shared_memory}: target ) &
    default( target map(mapper(vec_map),tofrom: vec) )
!$omp teams distribute simd
do i=1, vec%size()
    call vec(i)%work()
end do
!$omp end teams distribute simd
!$omp end metadirective
```

Nothing Directive

The nothing Directive

- The `nothing` directive makes meta programming a bit clearer and more flexible.
- If a certain criterion matches, the `nothing` directive can stand to indicate that no (other) OpenMP directive should be used.
 - The `nothing` directive is implicitly added if no condition matches

```
!$omp begin metadirective &  
    when( implementation={unified_shared_memory}: &  
          target teams distribute parallel do simd) &  
    default( nothing )  
do i=1, vec%size()  
    call vec(i)%work()  
end do  
!$omp end metadirective
```

Error Directive

Error Directive Syntax

■ Syntax (C/C++)

```
#pragma omp error [clause[[, clause],...]  
for-loops
```

■ Syntax (Fortran)

```
!$omp error [clause[[, clause],...]  
do-loops  
[!$omp end loop]
```

■ Clauses

one of: `at (compilation)`, `at (runtime)`

one of: `severity (fatal)`, `severity (warning)`

`message (msg-string)`

Error Directive

- Can be used to issue a warning or an error at compile time and runtime.
- Consider this a “directive version” of `assert()`, but with a bit more flexibility.

```
#pragma omp parallel
{
    if (omp_get_num_threads() % 2) {
#pragma omp error at(runtime) severity(warning) \
        message("Running on odd number of threads\n");
    }
    do_stuff_that_works_best_with_even_thread_count();
}
```

Error Directive

- Can be used to issue a warning or an error at compile time and runtime.
- Consider this a “directive version” of `assert()`, but with a bit more flexibility.
- More useful in combination with OpenMP metadirective

```
!$omp begin metadirective &  
    when( arch={fancy_processor}: parallel ) &  
    default( error severity(fatal) at(compilation) &  
            message(“No implementation available” )  
    call fancy_impl_for_fancy_processor()  
!$omp end metadirective
```

Programming OpenMP

OpenMP and MPI

Christian Terboven

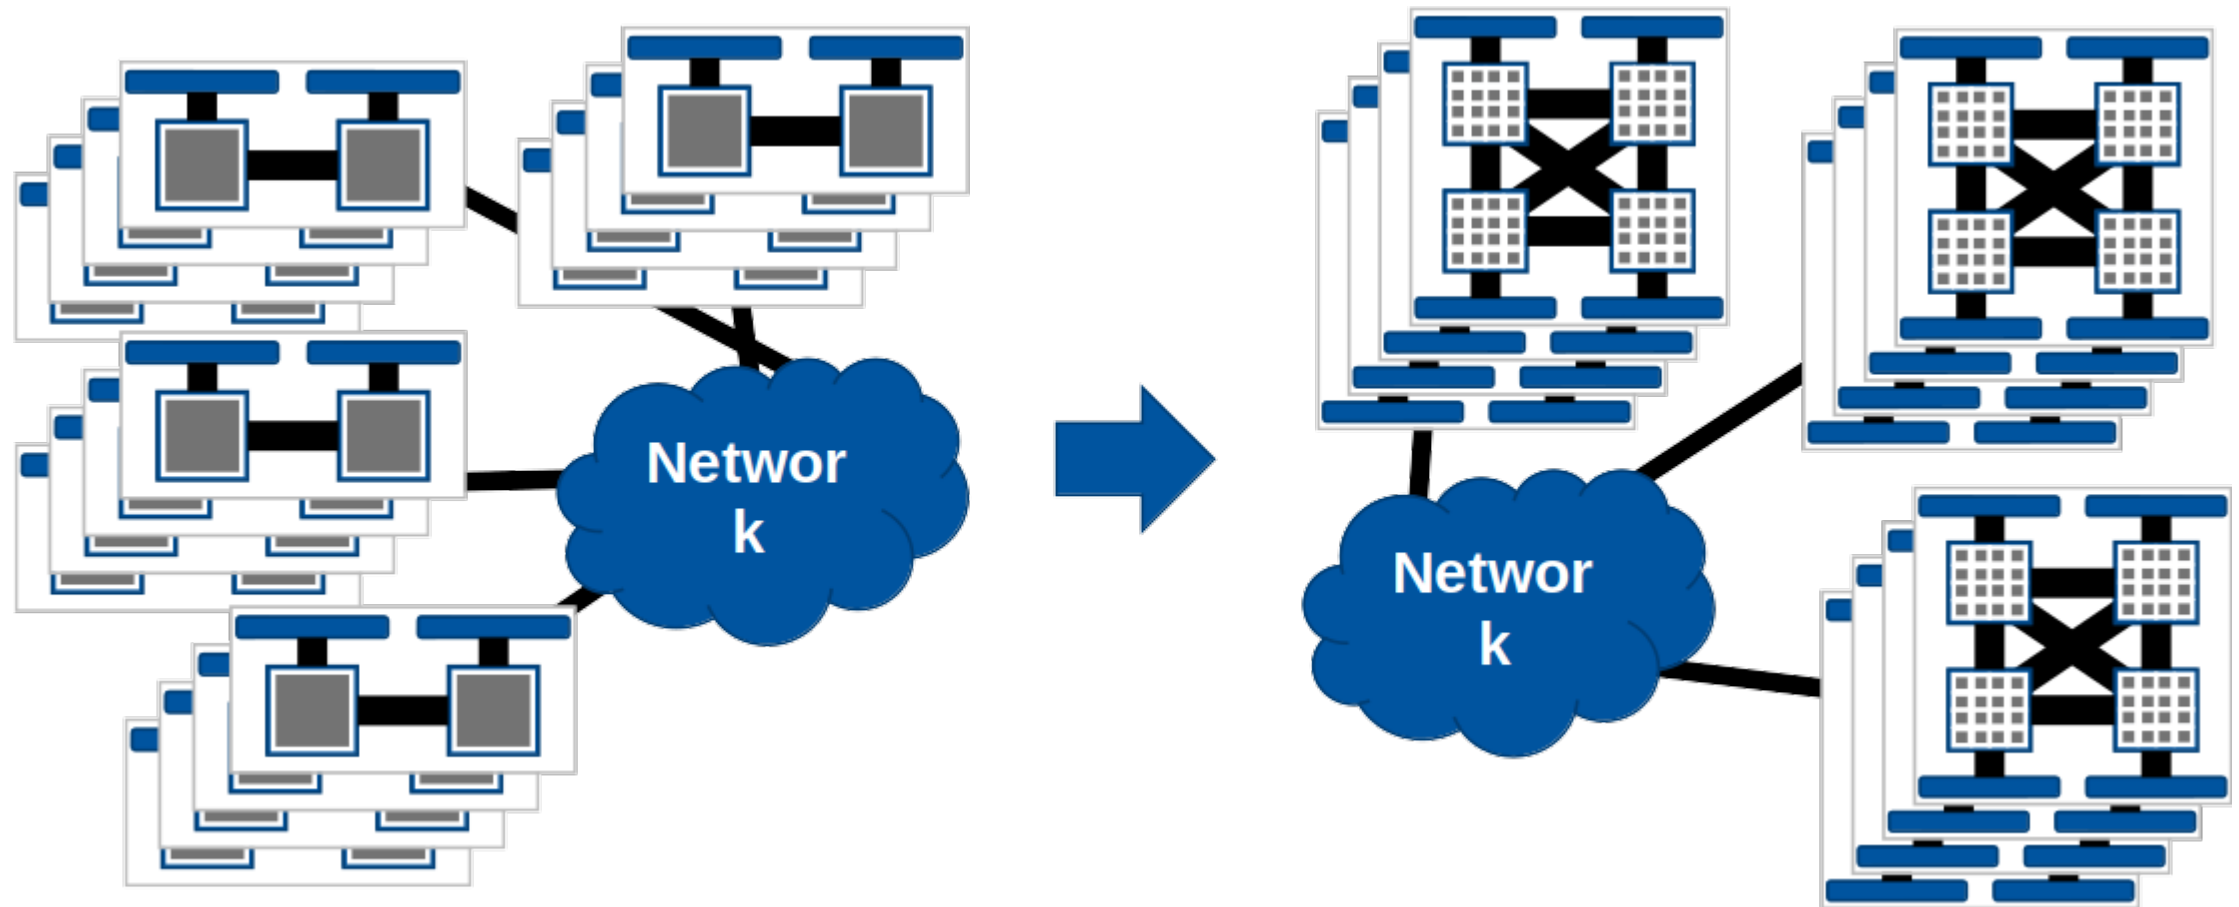
Michael Klemm



Motivation

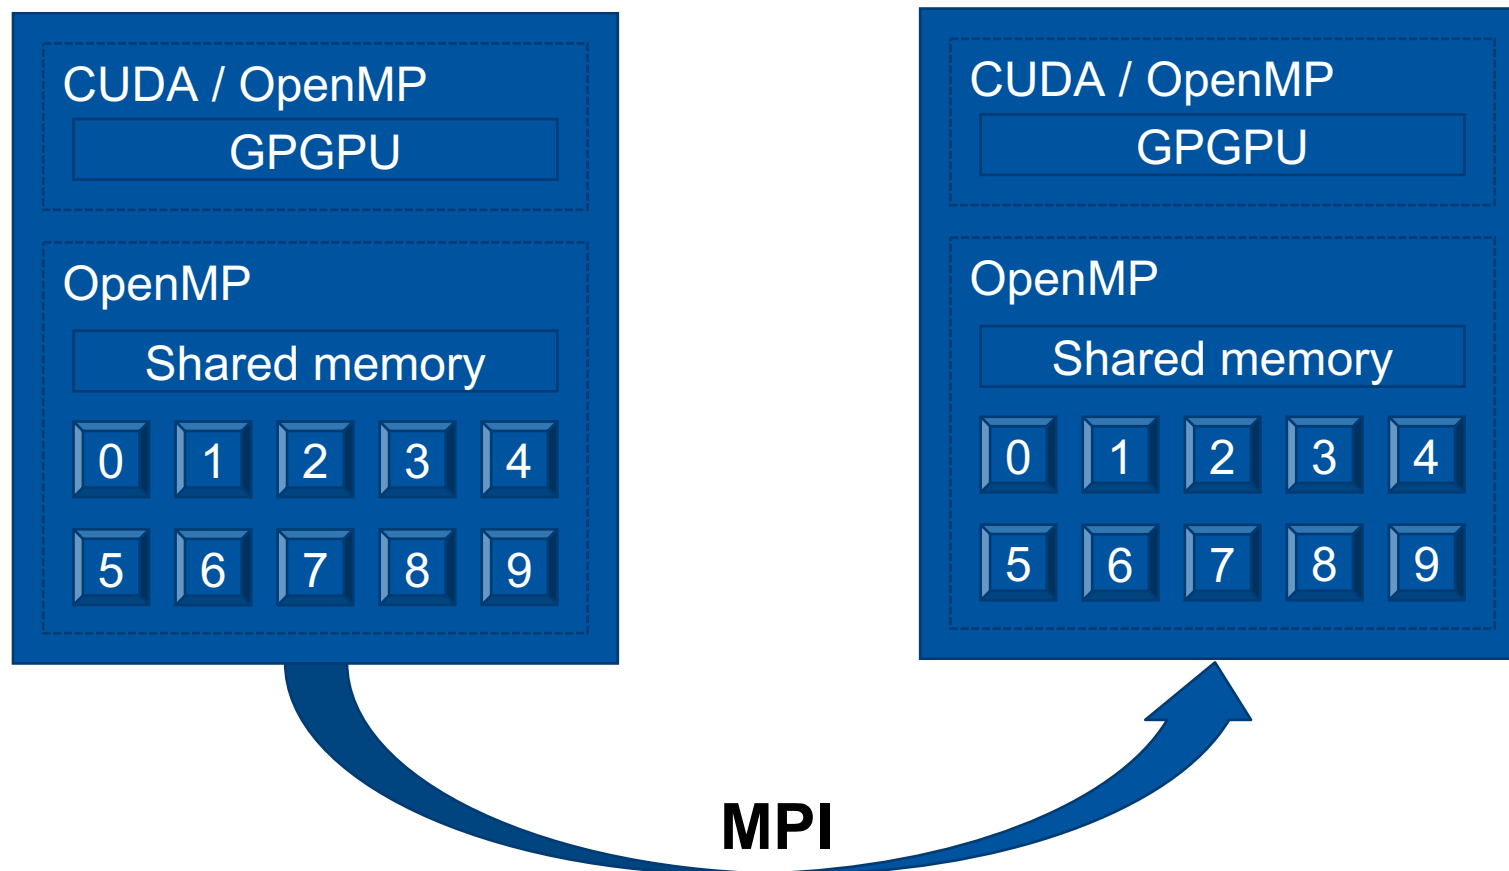
Motivation for hybrid programming

- Increasing number of cores per node



Hybrid programming


- (Hierarchical) mixing of different programming paradigms



MPI and OpenMP

MPI – threads interaction

- MPI needs special initialization in a threaded environment
 - Use `MPI_Init_thread` to communicate thread support level
- Four levels of threading support



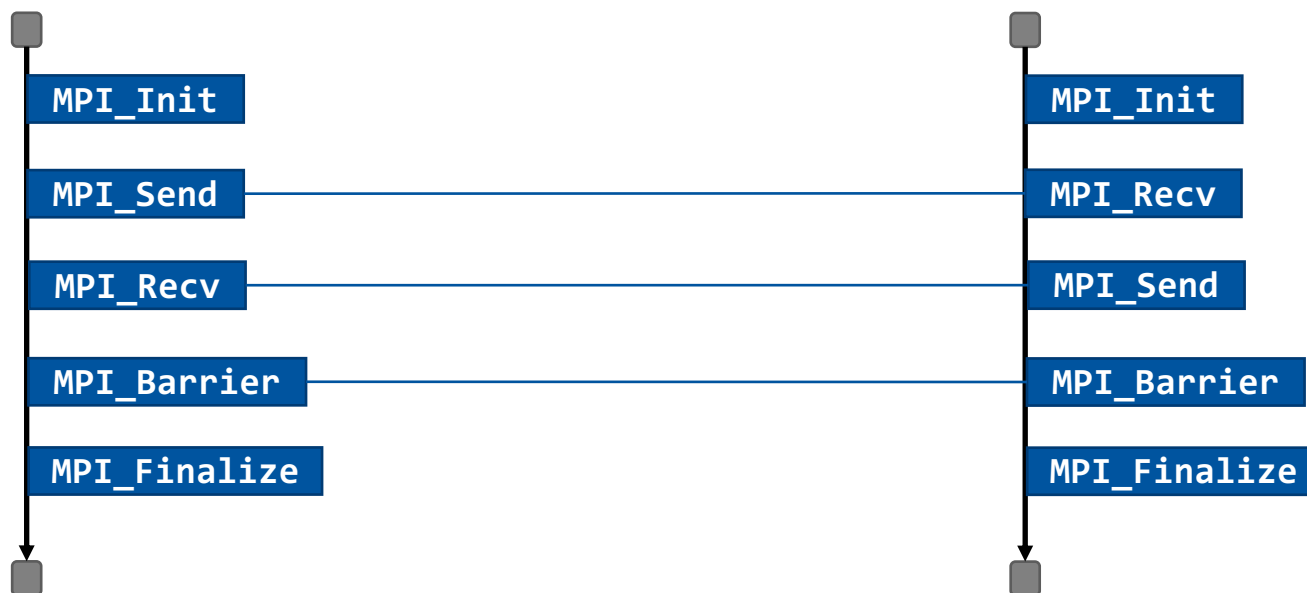
Level identifier	Description
<code>MPI_THREAD_SINGLE</code>	Only one thread may execute
<code>MPI_THREAD_FUNNELED</code>	Only the main thread may make MPI calls
<code>MPI_THREAD_SERIALIZED</code>	Any one thread may make MPI calls at a time
<code>MPI_THREAD_MULTIPLE</code>	Multiple threads may call MPI concurrently with no restrictions

- `MPI_THREAD_MULTIPLE` may incur significant overhead inside an MPI implementation

MPI – Threading support levels

- MPI_THREAD_SINGLE
 - Only one thread per MPI rank

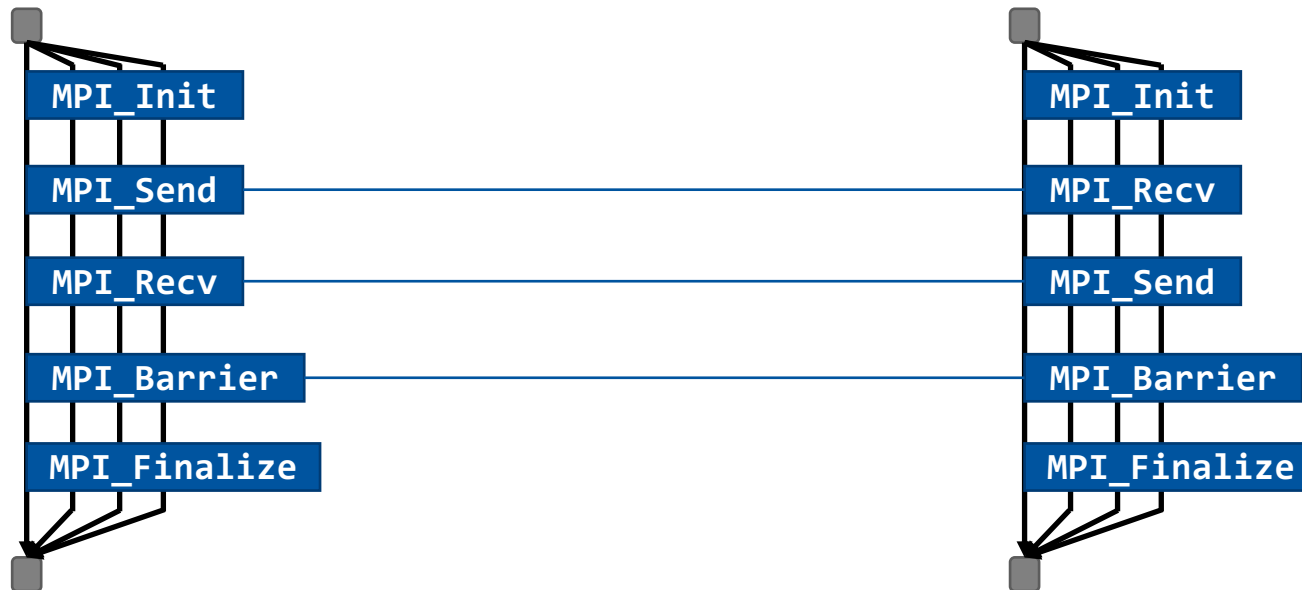
— MPI Communication
— Thread Synchronization



MPI – Threading support levels

- MPI_THREAD_FUNNELED
 - Only one thread communicates

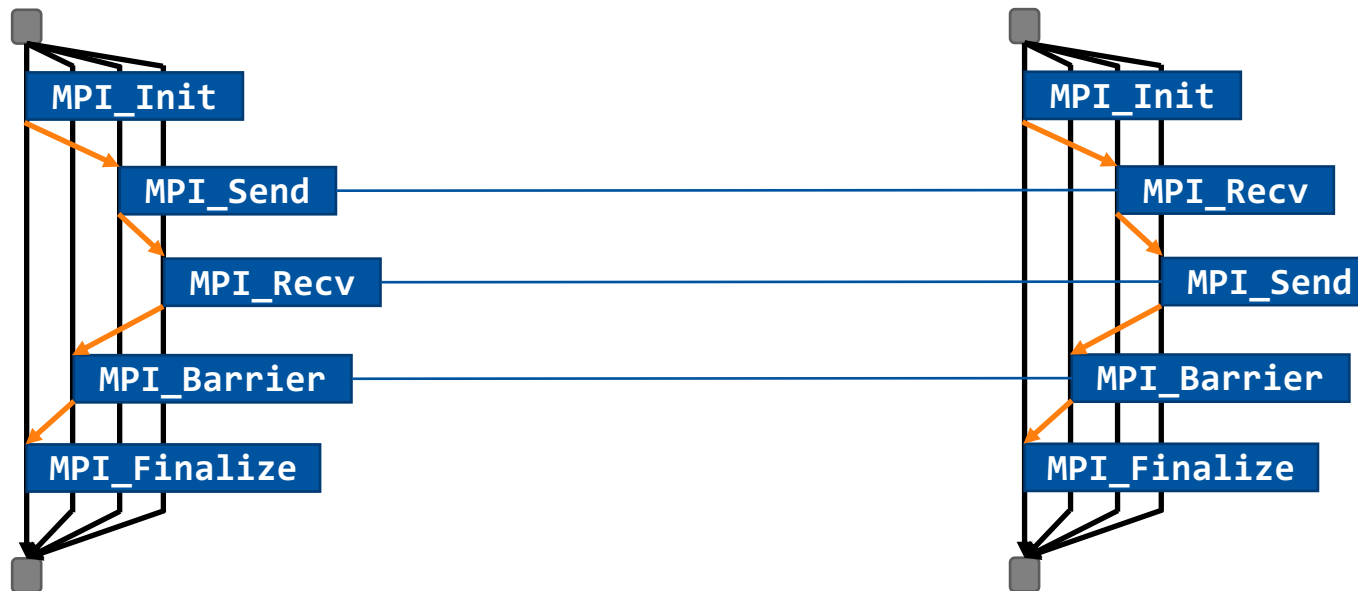
— MPI Communication
— Thread Synchronization



MPI – Threading support levels

- MPI_THREAD_SERIALIZED
 - Only one thread communicates at a time

— MPI Communication
— Thread Synchronization



MPI – Threading support levels

- MPI_THREAD_MULTIPLE
 - All threads communicate concurrently without synchronization

— MPI Communication
 — Thread Synchronization

